

GRAPH RAG, WIKIPEDIA / DBPEDIA, PUBLIC ONTOLOGIES, AND IEML

- **Date:** 2026-08-18
- **Updated:** 2026-08-18 — FR/EN/pt-BR dumps; Wikipedia as offline third-party source (not Arango); glossary for LOD and QID; multi-provider LLM handoff vs implicit semantics; schema.org and neighbouring public ontologies (names vs runtime)
- **Status:** idea / architecture note (not a spec, not an implementation plan)
- **Scope:** what Graph RAG is and why it exists; whether a local Wikipedia + DBpedia copy in French, English, and Portuguese is worth it as an offline source; whether schema.org / Dublin Core / SKOS / PROV-O / OWL add anything as ontologies; whether IEML adds anything beside graphs, dumps, and LLMs; mermaid translations of the 2009–2014 entity diagrams
- **Related:**
 - [14-proposed-architecture.md](#) — ASC control plane, conceptual graph, LOD / performance governor
 - [17-local-dev-stack-architecture.md](#) — extract-once, Solr + pgvector + Arango, paging, cost cliffs
 - [17-ui-design-ideas.md](#) — IEML as metaphor, not a hash codec; graph is conceptual
 - ASC: [Projet Complexe 2026 Revival \(v2\)](#)
 - ASC: [Would IEML really add tangible value for agents](#)
 - ASC: [Reasoning without Probabilistic Inference](#)
 - Origins: [Foundations of Task-oriented & Knowledge-oriented side-projects - v03.docx](#)
 - IEML corpus: Lévy 2023 HAL [hal-04055239](#); [sphere1.pdf](#); [00-0-sphere3-fr.pdf](#); [00-grammaire-ieml1.pdf](#); [IEML-Dictionnaire.pdf](#)

GOAL OF THIS NOTE


The 2026 revival already decided three things that this note must not undo:

1. **The graph is conceptual**, not necessarily one database (revival v2 §12). Solr, Postgres, Arango, files, and events can all *project* it.
2. **Knowledge is not merely RAG** (revival v2 §26). Claims, sources, evidence, unknowns, and knowledge-gaps are the model. Indexes are mechanisms.
3. **IEML is a compass, not a runtime** ([17-ui-design-ideas.md](#): do not encode IEML morphemes in the location hash).

The remaining question is how four *external* knowledge technologies sit on that architecture:

Technology	What it is, in one line	Layer it would occupy
Graph RAG	Build a graph from <i>your</i> texts, then retrieve a <i>bounded subgraph</i> (or a community summary) instead of the nearest embedding chunks	Indexer / query strategy over the personal corpus
Wikipedia articles + DBpedia triples (fr, en, pt)	A frozen public encyclopedia on disk: prose and/or RDF, consulted when the home link is down	Optional <i>offline third-party source</i> , not the second brain and not an Arango load
schema.org and neighbours	Shared <i>type and property names</i> (Person, Place, <code>dc:creator</code> , <code>wasGeneratedBy</code> , ...) used by search engines and Linked Open Data	Optional <i>spelling</i> of a few public-world types; not a graph to load and not the personal vocabulary
IEML	A constructed language whose <i>form</i> is meant to be semantically computable (syntagmatic <i>and</i> paradigmatic)	Optional <i>coordinate system</i> for concepts, if ever derived

They are complementary only if each one is used for the job the others cannot do. Mixing them into one “knowledge graph product” is how the project would freeze a bad ontology. schema.org is the most tempting freeze: it looks official, it already appears as `Thing (schema.org ?)` on the 2010s Histevents diagram, and it still does not name a KnowledgeGap.



Rule of thumb: Graph RAG over *your* notes. Wikipedia/DBpedia as a **local offline encyclopedia** (files you query, like opening a book when the internet drops) — not as nodes imported into Arango. schema.org as a **handful of type names**, not as an OWL stack. IEML as a *coordinate*, not as the language agents think in.

Terms used below

Two abbreviations in this note are easy to mix up with neighbouring jargon.
LOD here means **level of detail**, not “Linked Open Data”.
 In 3D and maps, you do not draw every brick of a city when you are looking at a country: you draw *less* as you zoom out. The proposed architecture uses the same idea for the graph on screen (and for how much you fetch):

Level	What you see / retrieve
LOD 0	Clusters / community summaries (“what is this region about?”)
LOD 1	Nodes
LOD 2	Relations
LOD 3	Labels
LOD 4	Metadata, citations, full passages

The *performance governor* picks a level from CPU/GPU/memory budget. “Render LOD” is that knob. (In Semantic Web writing, **LOD** often means *Linked Open Data* — DBpedia, Wikidata, RDF on the public web. This note spells that out as “Linked Open Data” when it comes up.)
QID is a **Wikidata item id**: the letter **Q** plus a number. Example: **Q1290** is Pierre Lévy. Same person, same QID, whether the Wikipedia article is in French, English, or Portuguese. A **P-id** is a *property* (**P31** = “instance of”). The personal graph can store a QID as a *pointer* (“this Author is that Wikidata item”) without loading Wikidata or Wikipedia into Arango.

1. GRAPH RAG

1.1 What it is

RAG (retrieval-augmented generation) means: before the model answers, fetch relevant text and put it in the prompt. Classic RAG embeds passages, retrieves the nearest neighbours, and generates from those chunks.

Graph RAG means: first turn a corpus into an explicit graph (entities, relations, optionally claims and communities), then retrieve a *structured neighbourhood* or a *precomputed summary of a cluster*, not only similar paragraphs.

It is still retrieval. It is not a reasoning engine, not ASC, and not “the knowledge model”. It is a way to choose *which* evidence to show an LLM when the question is about *connections* rather than *keywords*.

Microsoft Research’s pipeline (the one that named the genre) stores a small knowledge model:

Artifact	Role
Document	Input file or row
TextUnit	Chunk (default ~1200 tokens) with provenance
Entity	Person, place, work, event, ... extracted from a TextUnit
Relationship	Edge between two entities, with a description
Covariate (optional)	Time-bound claim about an entity
Community	Hierarchical cluster of entities (Leiden)
Community report	LLM summary of a cluster, used for <i>global</i> questions

That last pair is the distinctive move. Vector RAG has no object that answers “what is this corpus *about*, as a whole?”. Community reports do.

1.2 Why it exists

Vector RAG fails in a few recurring ways that this project already cares about:


Failure of chunk RAG	What actually happens	What Graph RAG tries
Global questions	“What are the main tensions in this archive?” retrieves random similar paragraphs	Map-reduce over community reports
Multi-hop	The answer is in A <i>linked to</i> B; A and B never co-occur in one chunk	Walk edges, or Personalized PageRank from seed entities
Entity identity	“Monnin” in three notes is three embedding islands	Merge entities; one node, many TextUnits
Contradiction / variant	Similarity retrieves <i>near</i> text, not <i>opposing</i> claims	Typed relation (contradicts , variant-of) if you extract it
Provenance	A chunk is a bag of tokens	Entity → TextUnit → Document is an addressable trail
Zoom	One retrieval granularity	Local neighbourhood vs community vs whole-graph summary

The 2010s Projet Complexe already had this problem without the name. The knowledge-oriented diagram’s **Link (pattern building)** is “I observed the same idea in several sources and previous notes”. Graph RAG is an *automated, lossy, expensive* attempt to propose those links. The human (or a later agent) still has to accept, type, and version them.

It also exists because **context windows are finite**. You cannot dump 55k files into a prompt. You need a *budgeted* view — which is the same constraint as the performance governor and IPC paging: never send the whole graph.

1.3 How it is implemented (the common pipeline)

All serious implementations share the same skeleton. They differ in *how expensive* extraction is, *how* they rank at query time, and *whether* they precompute community summaries.



Indexing is the cost cliff (same family as “ASR on every video” in [17-local-dev-stack-architecture.md](#)). Entity extraction is typically most of the LLM spend. Incremental updates are the hard part: Microsoft GraphRAG historically wanted a rebuild; LightRAG and temporal stores exist because corpora *change*.

Query must stay inside the same budgets as everything else here:

- **Query paging:** neighbours of one node, or one community level, never “all edges”.
- **IPC paging:** bounded JSON.
- **Render level of detail (LOD):** community report $\approx$ LOD 0 (clusters); entities $\approx$ LOD 1; relations $\approx$ LOD 2; chunk citations $\approx$ LOD 3–4. Same zoom idea as the graph on screen: do not fetch or draw everything at once.

Microsoft’s **local vs global** search is almost a restatement of knowledge-mode zoom:

Graph RAG search	Rough equivalent in this UI
Local	Focus on an entity; trail = neighbours; depth = hops
Global	Zoom out: cluster / century / domain; read a summary, not every edge
DRIFT (later MS variant)	Start global, spawn local questions, re-rank

That is useful. It is not a reason to vendor-lock the UI to Microsoft’s tables.

1.4 Leading open-source implementations (2026)

Brief list, not an evaluation harness. Prefer composing *behind a pivot* ([relate](#) / [research](#)) over picking a forever-stack.

- Task: alternative, quality / downside, requirement AND/OR/fallback
- Revival: `X is suspected, X conflicts with Z, X is sufficient for decision D`

If Graph RAG is used, **constrain extraction** to that closed vocabulary (schema-guided), or you will drown in mushy edges. “Schema-guided” here means *your* Type of Link / Histevents kinds — not “download schema.org and extract against 800 classes” (\$2.5).

1.6 Verdict

Use Graph RAG as **one retrieval strategy** over *selected* personal corpora (notes, papers, dissertations), behind `relate` / `research`.
Do not:

- run it on Wikipedia,
- run it on every social-media caption,
- let LightRAG/MS GraphRAG own identity,
- skip Solr (quotes, names, and filenames still win often).

Default order: **lexical** → **optional vectors on chosen chunks** → **graph walk on accepted entities**. Graph RAG is the third rung, for questions that are about *structure*.

2. LOCAL WIKIPEDIA + DBPEDIA AS AN OFFLINE THIRD-PARTY SOURCE (FR, EN, PT)

The point of a local copy is **not** to pour Wikipedia into Arango (or into the personal graph). It is to have a **third-party encyclopedia on disk** when the home connection is unstable: agents and the UI can still *look up* an article or a typed fact without wikipedia.org, dbpedia.org, or Wikidata being reachable. Arango stays the home of *your* relations (notes, tasks, claims, accepted links). Wikipedia stays a **separate library**. At most, a personal node stores a *pointer* (language + article title, or a Wikidata QID) so a lookup knows which file to open.
Two different dump families people glue together:

Dump	What it is	Order of magnitude (2026, compressed)
Wikipedia <code>pages-articles</code> per language	Current article wikitext (not full edit history)	enwiki ~25 GB; frwiki ~6.5 GB; ptwiki ~2.6 GB. Together ~34 GB bz2. Uncompressed XML is much larger.
DBpedia per language chapter	RDF extracted mainly from that language’s infoboxes + mappings to the DBpedia Ontology (DBO)	English-centric core historically ~0.85–1 billion triples. FR and PT chapters are extra, still far below the ~20 billion triples of <i>all</i> languages.
Wikidata <code>latest-all.json</code> (comparison, not required)	Live community knowledge graph, language-agnostic	~95–145 GB compressed JSON; uncompressed ~1+ TB. Heavier. Labels exist in fr/en/pt on the <i>same</i> QID.
Kiwix / ZIM (often the right <i>reader</i> format)	Offline snapshot meant to be <i>opened</i> , not parsed as XML	Typical Wikipedia ZIMs are smaller than raw XML; built for unstable links. Articles only, not DBpedia triples.

Wikimedia does not ship a separate “Brazilian Portuguese Wikipedia”. `ptwiki` **is Portuguese Wikipedia**, the edition used in Brazil (Brazilian Portuguese is a variety inside that wiki, not a second dump). `frwiki` and `enwiki` are the French and English encyclopedias.
Those three languages are the ones this project should cover if a dump is kept at all:

Language	Role in this life / corpus
French	Native. Theoretical and IEML/Morin/Lévy material.
English	Studied at the University of Sunderland (North East England), 2004–2008. Strongest overlap with most local LLMs and with DBpedia’s historical core.
Portuguese (Brazilian)	Living in Brazil since 2014. <code>ptwiki</code> for that language edition.

Three language editions are still in the same *order of magnitude* as the mixed personal archive in the 17-08 note (~85 GB / ~55k files). That is a different decision from “English only”, which would have cut the native and Brazilian sides of the knowledge-oriented work.

2.1 Advantages

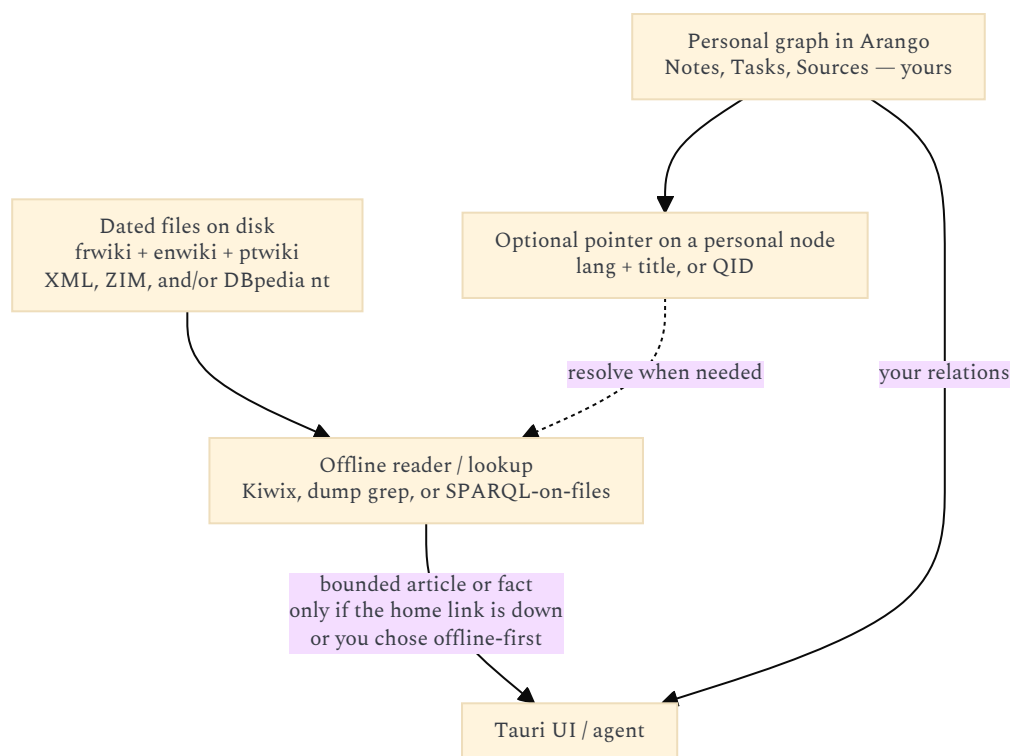
Advantage	Why it matters here
Works when the home link drops	The actual reason to keep a dump: lookup does not depend on wikipedia.org, SPARQL endpoints, or a VPN. Same idea as having PDFs on disk rather than “I’ll just open the URL”.
Stays out of Arango	The personal graph does not grow by millions of encyclopedia nodes. Wikipedia is consulted <i>like a book on the shelf</i> , not merged into Notes/Tasks.
Grounding / anti-hallucination	Agents can quote a <i>local</i> article (fr, en, or pt) or a <i>typed triple</i> (<code>dbo:birthPlace</code> , <code>rdf:type</code> <code>dbo:Philosopher</code>) instead of inventing a biography.

Advantage	Why it matters here
Stable snapshot	A dated dump is a reproducible world-state. Live websites move. Matches “what did we read while offline in 2026-08?”.
Pointer, not import	Store <code>frwiki:Pierre_Lévy</code> / <code>enwiki:Pierre_Lévy</code> / <code>ptwiki:Pierre_Lévy</code> or one QID on <i>your</i> Author node. Resolve against the local dump at query time. No need to ingest the article body into Arango.
Three working languages	A French note can hit <code>frwiki</code> ; an English paper, <code>enwiki</code> ; a Brazilian source, <code>ptwiki</code> . Same person can still share a QID across the three.
Schema.org / “Thing” already in Histevents	The 2010s diagram pointed Location/Group at Thing (schema.org ?) . That “?” is answered in §2.5: DBpedia/Wikidata are the <i>instance</i> library; schema.org is only a <i>class spelling</i> . Neither is your entity table.
English still helps models	Many local LLMs are strongest in English. You can still <i>prefer</i> <code>enwiki</code> for model-facing lookups without deleting <code>frwiki</code> / <code>ptwiki</code> for you.
License is known	Wikipedia text is CC BY-SA. DBpedia inherits that world. Better than a random crawl. Share-alike applies if you <i>publish</i> derivatives.
Cheap relative to Graph-RAG-ing Wikipedia	Keeping files (or a Kiwix reader, or a small SPARQL over dump files) is finite disk. LLM-extracting millions of articles is the cost cliff you must refuse.

2.2 Inconvenients

Inconvenient	Why it hurts this project
Disk, RAM, backup	Three article dumps (~34 GB compressed) plus optional triples compete with PDFs, photos, and video. Indexes or ZIM readers add more. Feasible on dedi HDDs; tight on a seven-year-old laptop.
Staleness	Wikipedia changes daily. A local copy is only as fresh as the last download. Agents will be <i>wrong in a dated way</i> while offline. Version the snapshot (<code>dump--frwiki-2026-06-01</code> , etc.).
Infobox noise (if you keep DBpedia)	DBpedia is a parse of messy infoboxes. Types are incomplete; predicates collide; long-tail articles are empty. Treat as <i>candidates</i> , not ground truth.
Three editions ≠ three identical worlds	The French, English, and Portuguese articles on the same topic disagree, omit, and bias differently. Offline lookup must say <i>which language edition</i> was used. A QID identifies the <i>item</i> ; it does not make the three texts the same.
<code>ptwiki</code> is not “Brazil-only”	Portuguese Wikipedia mixes European and Brazilian usage and topics. Good enough as the pt edition; not a national encyclopedia of Brazil.
Wrong layer if it leaks into the default graph	Revival §26: knowledge is claims, unknowns, gaps. Encyclopedias are <i>sources</i> . If Wikipedia titles start appearing as first-class graph nodes, the UI becomes an offline Wikipedia with a sidebar. Pointers only.
Graph RAG on Wikipedia is forbidden-cost	Community detection + LLM summaries over even one language edition is an institutional budget. Use the dump as <i>text or triples to read</i> , do not re-extract it.
Query shape (triples)	SPARQL over hundreds of millions of triples needs a dedicated engine, paging, and a warm process. That engine is not Arango. Same “never dump the whole graph to the webview” rule.
DBpedia vs Wikidata	Wikidata is the living public KG (one QID, labels in fr/en/pt). DBpedia is the <i>Wikipedia-shaped</i> extract per language. For “typed facts while offline”, a Wikidata truthy subset or Kiwix-style articles may beat three DBpedia chapters. Decide explicitly.
License share-alike	Publishing a site/ebook that embeds Wikipedia text pulls CC BY-SA onto those pages. The <code>publish</code> pivot must know that.
Identity mess if you import anyway	Three titles + DBpedia URIs + QID + your Author node. Without a pointer policy, Graph RAG extraction will mint a fifth. Importing into Arango makes this worse; keeping dumps separate avoids most of it.
Not multimodal	Dump is text/RDF. It does not help photos, ASR, or sheet music. The heterogeneous pipeline stays.
Legal/ethical scrape hygiene	Use official dumps (or Kiwix ZIMs), not a crawler. Large third-party corpus: keep it out of git.

2.3 How it should sit (if at all)



The reader is an ASC-supervised **lookup**, like opening a file. It is not a fan-out indexer that copies Wikipedia into Solr+Arango by default. (You *may* later put *selected* articles into Solr if you search them often. That is still a copy of passages you chose, not “Wikipedia lives in Arango”.)

Strategy	Verdict
Load Wikipedia/DBpedia into Arango as graph nodes	No. That was never the idea. Swallows the personal graph; fights the “conceptual graph” rule.
Kiwix / ZIM for fr + en + pt	Strong default for articles while offline. Human-readable; agents can be given a CLI lookup. No triples.
Raw XML dumps on disk + a small extractor at query time	Fine if you want scripts (<code>get abstract for title X in lang L</code>). Heavier than ZIM. Still not Arango.
DBpedia nt files + a dedicated SPARQL (or even <code>grep</code>) beside Compose	Only if you need <i>typed facts</i> offline. Keep that store next to Arango, not inside it.
Live wikipedia.org / SPARQL when the link is up	Fine. The dump is the <i>fallback</i> , not the only path.
Watchlist of titles/QIDs that appear in your notes	Optional cache: copy <i>those</i> articles into the offline reader first. Much smaller than three full encyclopedias.
Wikidata truthy subset instead of three DBpedia chapters	Often simpler for “same item, three labels”. Still a separate file/engine, not Arango.
Graph RAG over the Wikipedia dumps	No.

2.4 Verdict

Keep **French, English, and Portuguese** snapshots as an **offline third-party library** (Kiwix/ZIM and/or dump files; DBpedia or a Wikidata subset only if typed facts matter while disconnected).

Do **not** ingest them into Arango. The personal graph may store **pointers** (titles per language, or a QID). The default knowledge view stays *your* notes, claims, and gaps.

Prefer **online lookup when the home connection works**, local dump when it does not. Never Graph-RAG the encyclopedia.

2.5 SCHEMA.ORG AND OTHER PUBLIC ONTOLOGIES

The Histevents diagram already asked this: Location and Group point at **Thing (schema.org ?)**. That is a *vocabulary* question, not a dump question. schema.org does not ship Pierre Lévy. It ships the class `Person`. Wikidata `Q1290` and `frwiki:Pierre_Lévy` are instances. Mixing those two jobs is how “we should use schema.org” quietly becomes a second Drupal taxonomy.

2.5.1 What these things are

Public ontologies are **shared type and property names**. They answer “what *kind* of thing is this?” in a language search engines and Linked Open Data already use. They are not a second encyclopedia, not Graph RAG, and not IEML.

The widely used ones, for this project’s purpose (not an ontology census):

Vocabulary	What it actually is	Typical classes / properties	Relation to this stack
schema.org	Vocabulary for marking up web pages so Google / Bing / Yandex can make rich results (usually JSON-LD). The most <i>deployed</i> “ontology” on the public web.	Thing → Person, Organization, Place, Event, CreativeWork, Product, ... plus thousands of properties	The Histevents “?” — public- <i>world</i> types
Wikidata (P31 and properties)	Living community knowledge graph; language-agnostic items	instance of (P31), occupation (P106), ...	Already the <i>pointer</i> layer (QID). Classes are <i>your</i> P31, not schema.org imports
DBpedia Ontology (DBO)	Types extracted from Wikipedia infobox mappings	dbo:Philosopher, dbo:birthPlace	Comes with a DBpedia dump; English-centric. Already in §2
Dublin Core (DC / DCTERMS)	Tiny bibliographic set, 1990s–now; still what librarians and <i>meta</i> tags use	title, creator, date, rights, type	Names for Source / Assembly fields. Solr already thinks this way
SKOS	Thesaurus pattern (W3C)	broader / narrower / prefLabel / altLabel	Same <i>shape</i> as Topic / Domain <i>parent</i> loops
FOAF	People and homepages (2000s Semantic Web)	Person, knows, mbox	Largely absorbed by schema.org <i>Person</i> + Wikidata. Ignore as a stack
PROV-O	W3C provenance	wasGeneratedBy, used, wasDerivedFrom	Same <i>job</i> as revival §27 (<i>extracted_by</i> , provider, time)
RDFS / OWL	The <i>meta</i> languages for writing ontologies	subClassOf, equivalentClass, reasoners	How DBpedia maps dbo:Person to schema:Person. Not a domain to implement
Open Graph Protocol	Facebook/LinkedIn link-preview markup	og:title, og:type	Even thinner web-page metadata. Irrelevant inside Tauri

Honourable mentions that look serious and would be a mistake here: **CIDOC-CRM** (museums and cultural heritage), **BFO** / **DOLCE** / **SUMO** / **Cyc** (upper ontologies), **BIBO** / **FRBR** (library theory). Those are careers, not Compose pivots. DBpedia already declares things like `dbo:Person owl:equivalentClass schema:Person`. Wikidata items often carry the same alignment. If you keep dump files, you **already have** schema.org alignment without adopting schema.org as a product.


2.5.2 Would they add value?

Almost none as a runtime. A little as stolen names.

The project already has three type systems that matter:

1. **Personal interpretive vocabulary** (knowledge v08 / task v09): Note, Claim, Link, Factor, KnowledgeGap, Requirement, Fallback — *not* in schema.org.
2. **Public identity** (QID / wiki title): *this* Person — not a class.
3. **Closed relation types** for Graph RAG: contradiction, complementarity, variant — schema.org’s `relatedTo` / `mentions` are mushier than v08.

schema.org’s *Thing* tree overlaps **Histevents** (Person, Place, Organization, Event, CreativeWork) and almost completely misses **knowledge-oriented** and **task-oriented** objects. Google never needed a KnowledgeGap.



Advantages (the thin slice)

Advantage	Why it might still matter
Shared names for public-world kinds	Saying <code>type: Person</code> (meaning schema.org/Person) on a node that also has a QID is a <i>convention</i> other tools understand. Cheap documentation.
Constrains Graph RAG extraction	A <i>tiny</i> allowlist (Person, Place, Organization, Event, CreativeWork) is better than free-form entity types. The 2010s Histevents types already <i>are</i> that allowlist; schema.org is one published spelling of it.
Dublin Core field names	<code>dc:title</code> / <code>dc:creator</code> / <code>dc:rights</code> for Source metadata. No reason to invent <code>titre_du_document</code> .
PROV-O names	<code>wasGeneratedBy</code> $\approx$ <code>extracted_by</code> + provider. Aligns provenance YAML with a known vocab without running a reasoner.
SKOS pattern	<code>broader is parent</code> on Topic/Domain. Keep the parent loop; do not stand up a SKOS triple store.
JSON-LD as an export later	If <code>publish</code> ever emits a public page, schema.org JSON-LD is how search engines read it. That is a <i>publish</i> concern, not the second brain.
Mappings you do not have to maintain	DBpedia $\leftrightarrow$ schema.org $\leftrightarrow$ Wikidata already exists. You do not need a local OWL reasoner to know Person is Person.

Inconvenients (why they mostly do not pay rent)

Inconvenient	Why it hurts this project
Wrong job	schema.org exists so a search engine can rich-result a recipe. This app is not crawled. There is no consumer of your JSON-LD except you.
Giant schema vs revival	schema.org has on the order of 800+ types and thousands of properties. Loading it as “the ontology” is the freeze-a-bad-schema failure already named in the intro.
Misses the objects you actually care about	No Claim, Evidence, KnowledgeGap, Type of Link (<code>contradiction</code>), Factor, Requirement, Fallback, Completion. Those are Projet Complexe’s value.
Types $\neq$ identity	<code>schema:Person</code> is a class. Pierre Lévy is a QID. The 2010s “?” was already answered by DBpedia/Wikidata as a <i>library</i> , not by importing schema.org classes into Arango.
Fourth (fifth) name for one idea	Personal slug + wiki title + QID + <code>dbo:</code> + <code>schema:Person</code> + IEML USL. Same identity mess as §2.2.
JSON-LD in the Tauri webview	A second ontology in the frontend, which 17-ui-design-ideas.md already forbids.
FOAF / OGP / GoodRelations	Absorbed, commercial, or dead. Ignore.
Upper ontologies (BFO, Cyc, SUMO, DOLCE)	Philosophers arguing about continuants. The 2012 Drupal-taxonomy trap with more axioms.
CIDOC-CRM / FRBR	Histevents- <i>adjacent</i> and still a research career. Not a first pivot.

Inconvenient	Why it hurts this project
OWL reasoning	Classification as a decision procedure is attractive (Lévy, Morin, “architecting my mind”) and is how you lose years. The performance governor and KnowledgeGap are the 2026 answer to overload, not <code>owl:disjointWith</code> .

2.5.3 How it should sit (if at all)

Strategy	Verdict
Import schema.org (or OWL, SKOS, CIDOC) into Arango as the type system	No. Same as “load DBpedia into Arango”.
Mark every Note as JSON-LD <code>schema:Article</code>	No. Notes are epistemic objects, not web pages.
Tiny allowlist of public-world types, spelled as Histevents names or as schema.org URLs	Optional. Prefer the 2010s names (<code>Person</code> , <code>Location</code> , <code>Group</code> , <code>Event</code>). If you want an external spelling, <code>schema:Person</code> on <i>those</i> nodes only, next to a QID.
Dublin Core keys on Source	Fine. Field names, not a graph.
PROV-O-shaped provenance on Claims	Fine. Names, not a reasoner.
SKOS for Topic.parent	The pattern already exists. Do not add a SKOS database.
schema.org JSON-LD at <code>publish</code> time	Later, if a public site exists. Not the local graph.
Align IEML USLs to schema.org classes	No. Two research ontologies glued together.

2.5.4 Verdict

schema.org does not add a new capability. Identity is QID / title pointers. Typed facts while offline are DBpedia or Wikidata dumps. Corpus structure is Graph RAG. Sense / paradigms is IEML-the-compass. Personal epistemology is v08 / v09.

What is worth stealing: a **dozen names** (`Person`, `Place`, `Event`, `CreativeWork`; `dc:creator`; `wasGeneratedBy`) so Graph RAG extraction and Source metadata stay boring. That is schema-guided extraction in §1.5, not an ontology project.

Using schema.org (or SKOS, FOAF, OWL, CIDOC) as a *system to implement* would add another ontology and still would not name a KnowledgeGap. The 2010s “Thing (schema.org ?)” is best read as: **public-world kinds may reuse a well-known spelling; the public instances stay in the offline library; your Claims stay yours.**

3. IEML AS A COMPLEMENTARY TOOL

Sources for this section: Lévy, *Calculer la sémantique avec IEML* (2023, HAL); *La sphère sémantique*; the IEML grammar and dictionary PDFs; the 2010s [IEML-Analysis-v07](#); and the 2026 memos on agents.

3.1 What Lévy is actually claiming

IEML (**I**nformation **E**conomy **M**eta**L**anguage) is a **constructed language** meant to have:

- the **expressive power of a natural language** (recursive phrases, roles, modalities, narrative),
- computable linguistic semantics**, including the part linguistics never fully mathematized: the **paradigmatic** axis (substitution, matrices of related concepts), not only the syntagmatic tree (phrase structure).

Concrete pieces (2023 paper):

Piece	Claim
Alphabet	Six primitives (classically <i>U, A, S, B, T, M</i> — void/virtual/actual, sign/being/thing; notations vary by document generation)
Dictionary	~3000 words, each in exactly one paradigm ; words defined by phrases using other words (circular interdefinition, as in natural language)
Phrases	A syntagmatic function with a fixed set of roles; phrases may belong to <i>several</i> paradigms, freely created
USL	<i>Uniform Semantic Locator</i> — canonical form, URI-compatible, unique
Graphs	Links woven from <i>syntactic conditions</i> , producing knowledge graphs / hypertexts
Literals	Proper names, URLs, GPS, dates sit in <code><references></code> — <i>not</i> IEML-computable; contrast with RDF URIs, which Lévy calls rigid designators whose <i>string form</i> carries no semantics
Editor	Open-source demonstrator: write with FR/EN words, parser produces IEML. Humans should not have to see the raw code

Lévy’s applications, in his order: (1) **metadata for digital memory** / collective intelligence, (2) **neurosymbolic AI** (IEML as protocol between humans, between humans and machines, and between machines), (3) **digital humanities** as a semantic coordinate system.

He is explicit that raw IEML is for machines; people see natural language, diagrams, icons. That already agrees with “do not put IEML in the hash”.

The older *Sphère sémantique* / grammar stack is heavier: rhizomes, circuits (paradigmatic, syntagmatic, chronotopic, hypertextual), *cités*, algebraic operations (rotation, powerset, partition). [IEML-Analysis-v07](#) is a 2010s attempt to *entity-diagram* that stack. It is a research map, not a Compose service.

3.2 Why this was attractive to Projet Complexe in the 2010s

The Foundations doc already names the wound Graph RAG still has not healed:

- knowledge is a **living process** (loops, decay when unattended),
- storing it forces **classification**, which forces **organization** (“architecting my mind”),
- **information overload**: relating everything becomes infinite digression — hence the need for focus (“just enough research”),
- opinions **change** → versioning,
- **naming things** (epistemology, topic maps, semantic web),
- unity of the two orientations: knowledge vs tasks are *complémentaire, concurrent et antagoniste* (Morin).

Lévy and Morin are listed there as the *corpus de référence*. IEML was the hope that classification could be **computable** rather than a pile of Drupal taxonomies.

DBpedia/RDF answered *reference* (“this URI denotes Rodin”). schema.org answered *kind for the web* (“Rodin is a **Person** / **CreativeWork** in Google’s dialect”).

IEML claims to answer *sense* (“sculptor” as a position in a paradigm, transformable by algebraic operations). Graph RAG answers *corpus structure* (“in my notes, Rodin is linked to these claims”). Four different questions. schema.org does not replace the other three; it barely overlaps Histevents.

3.3 The 2026 agent question: still mostly redundant, in a precise sense

The memo [Would IEML really add tangible value for agents](#) remains the right split:

LLMs already have a latent IEML — paraphrase, analogy, taxonomy induction, ontology mapping — without a constructed language.

The original 2010 premise (“machines have no semantics”) is false **for a single model, in a single sitting**. That model *does* have semantics. They are just trapped in its weights and in the current context window.

What remains is the difference between **implicit** and **explicit** semantics:

Implicit (LLM / embeddings)	Explicit (IEML / RDF / typed graph)
Cat ≈ feline ≈ pet	Inspectable path through intermediate concepts
Model-dependent	Can survive a model upgrade and a provider swap mid-task
No persistent object (Monnin)	Addressable USL / URI / entity id
Flow-level intervention (Meadows)	Information- <i>structure</i> intervention
Lives in one chat / one weights file	Can be read by the next model, or by inspect-agent

IEML-the-language is still unnecessary for one-shot summarization, coding, email, brainstorming, and RAG over a few thousand documents **when one model does the whole job**.

It (or, cheaper, a **typed graph with closed vocabularies**) becomes interesting as soon as work has to **survive a change of mind** — years later:

Is today’s concept X the same as what we called Y in 2016?

and also **minutes later**, when ASC swaps the engine under a stable pivot.


3.3.1 ASC swaps implementations and LLM providers

Revival v2 already states the consumer should ask for a **capability**, not an implementation (**pdf-totext** vs Docling vs Tika; local LLM vs remote API vs another runtime). Projet Complexe ASC is that thin name: **extract**, **research**, **run-agent** stay put while the hook behind them changes.

That list of LLM backends is not hypothetical. A single environment is expected to mix:

Provider	Where it runs	Typical use
Ollama	This laptop	Small / simple models; cheap, private, offline
Tiiny.ai (or similar)	A device on the LAN	A bit more elaborate local inference without leaving the house
Remote model APIs	Internet	Heavier reasoning when the home link and budget allow
Cursor CLI (wrapped)	Local agent runtime that itself calls models/tools	Coding, repo inspection, “an agent with files and tools” rather than a raw completion

ASC’s job is to make **run-agent** / **research** / **relate** **the pivot** and the provider a *sidecar choice* (Requirement / Environment / Technology in the old task diagrams: “needs GPU”, “must stay on LAN”, “may call a remote API”). The UI and the task object must not hard-code **ollama run ...**.



The **pitfall** is not swapping models between *different* tasks. It is one task **partially** handled by several of them: Ollama drafts entities, a remote API judges a contradiction, Cursor CLI patches a file, Tiiny resumes tomorrow. Each model has its own latent map. They do **not** share “cat ≈ feline”. They only share what you **wrote down** in a form the next caller can parse.

Pitfall if the only memory is implicit	What goes wrong mid-task	What explicit structure avoids
Entity drift	Model A extracts “Alexandre Monnin”; model B creates “Monnin” as a new node	One Author / QID / entity id; later models <i>resolve</i> , they do not mint
Relation-type drift	A says <code>contradicts</code> ; B writes “disagrees somewhat” as prose	Closed Type of Link (<code>contradiction</code> , <code>complement</code> , ...) from the knowledge diagram
Embedding mix	Vectors from Ollama’s embedder are not comparable to an API embedder; mixing them silently ruins RAG	Named embedding space on the vector projection; do not fan-out incompatible vectors into one index
Handoff amnesia	Model B never saw model A’s hidden state — only the last chat blob, if that	Task state, trail, KnowledgeGap, Completion as objects, not “whatever was in the prompt”
Confidence laundering	A 7B local guess becomes “fact” when a frontier model continues from it	Provenance: <code>extracted_by</code> , model id, provider, time, confidence — revival §27
Language / register shift	FR note → EN-centric API; pt-BR source → small Ollama; strings diverge	Same entity id; language is a <i>property of the Source</i> , not of the identity
Shape mismatch	Cursor CLI returns a diff; Ollama returns a paragraph; an API returns JSON	Pivot contract: bounded, typed artifacts (YAML/JSON sidecar), not free-form stdout as the system of record
Privacy / budget surprise	Step 4 of a private task silently hits a remote API	Requirement/Environment on the task: <code>lan-only</code> vs <code>api-ok</code> ; ASC enforces, the model does not decide
Resume after killswitch	Research paused; another model restarts and re-litigates settled claims	Explicit KnowledgeGap + “what is already sufficient for decision D”

So: **yes**, explicit semantics help avoid those pitfalls — but the urgent explicit layer is the **typed personal graph + pivot I/O schema + provider provenance**, not IEML morphemes.

- **ASC** names the capability and which provider ran (`run-agent` with `provider=ollama|tiiny|api|cursor-cli`).
- **Projet Complexe** owns the meaning that must survive the hop (Claim, Link, Factor, Task, Requirement, KnowledgeGap).
- **IEML** would only be an extra canonical *address* on durable Concepts if string ids and QIDs start to collide across years and languages. It is not required to make Ollama and Cursor CLI interoperable tomorrow.

The 2010 sentence is therefore better put like this:

Machines *do* have semantics — each machine, each session, each model. They do not have **the same** semantics, and they do not keep them. ASC’s reason to exist is to keep the **name of the operation** stable while implementations (Tika or Ollama or Cursor CLI) change. Explicit structure is what keeps the **name of the thing being operated on** stable across those changes.

Without that, swapping providers is only half of portability: you can change the engine and still **lose the work** at every handoff.

3.4 Advantages of IEML as *complement* (not as replacement)

Advantage	Complement to Graph RAG	Complement to Wikipedia/DBpedia
Canonical address in meaning-space (“compass”)	Graph RAG nodes are <i>strings extracted from your prose</i> (“resilience”). IEML would say <i>which</i> resilience.	DBpedia <code>dbr:Resilience</code> is a Wikipedia page, often a disambiguation mess. USL aims at the <i>concept</i> , not the article.
Paradigmatic computation	Community reports are statistical clusters. IEML paradigms are <i>generated matrices</i> (roles × variables). You can ask for symmetries, not just “related entities”.	RDF <code>skos:broader</code> is asserted. IEML claims some relations <i>follow from form</i> .
Cross-language	Extraction in FR vs EN vs PT yields different entity strings.	Three Wikipedia editions disagree in prose; a QID still names <i>one</i> item. IEML’s editor is FR/EN; Portuguese would be extra work.
Survives wording and model swaps	Re-extracting with a new LLM reshuffles the graph. Mid-task hops (Ollama → API → Cursor CLI) do the same unless the pivot writes typed objects. A stored USL / entity id need not reshuffle.	Wikipedia titles rename; redirects help, but the <i>sense</i> still sits in prose.
Interop protocol between agents and providers	Agents today pass JSON English (<code>{"goal":"find papers"}</code>). A semantic object can be handed to the next model without re-parsing a chat log.	Public KGs do not know your task vocabulary (Requirement, KnowledgeGap, <code>lan-only</code>).
Literals vs concepts	Graph RAG mixes “Rodin” and “sculptor”. Lévy splits them: name in <code><Rodin></code> , concept in the phrase.	DBpedia also splits (resource vs infobox text), but URIs do not encode the split in their <i>shape</i> .
Meadows primitives as a paradigm	You can Graph-RAG the word “stock”. You cannot <i>rotate</i> it toward “flow” as an algebraic operation unless the language supports it.	Wikipedia has articles; it does not have a generated systems-theory matrix.
Metadata ethics	Lévy: semantic metadata are political (who organizes digital memory). A closed IEML layer is at least <i>inspectable</i> .	GAFAM KGs and Wikidata are inspectable in theory, ungovernable in practice for a personal stack.

The durable part may not be the language. It is the **checklist** already isolated in the 2026 memo:

compositionality, computability, canonical representation, invertibility where possible, stable identifiers, semantic interoperability.

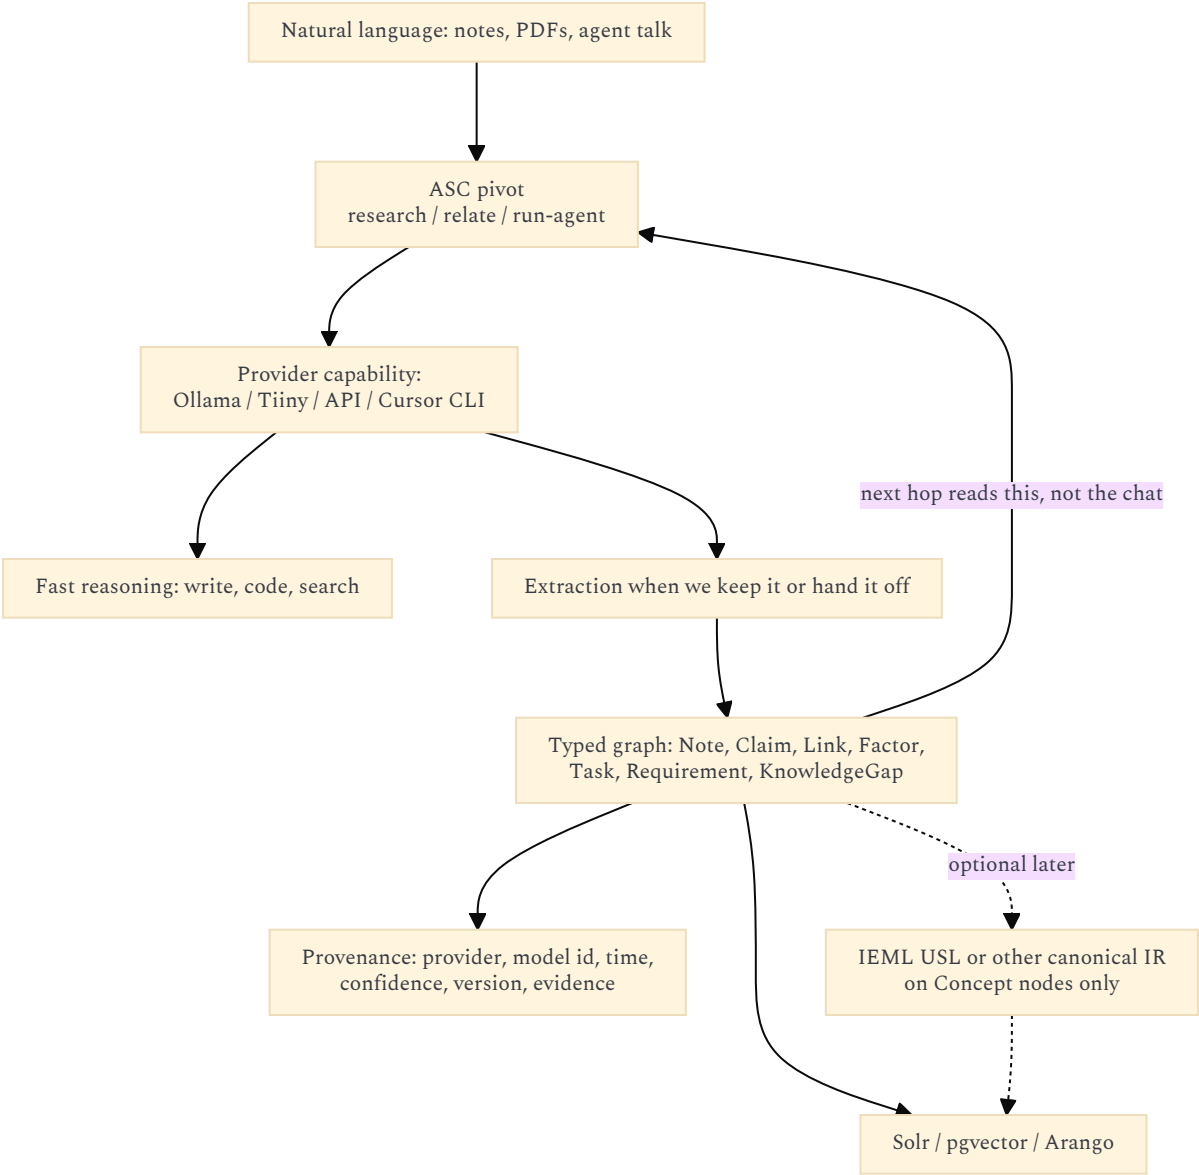
Those can be satisfied by RDF, property graphs, or a tiny typed IR. IEML is one proposed solution to that spec.

3.5 Inconvenients

Inconvenient	Consequence
Redundancy with LLMs for most <i>single-model</i> work	One Ollama call to summarize a note still does not need IEML. Multi-provider <i>handoff</i> needs a typed IR; that IR can be YAML/Arango, not IEML.
Tiny ecosystem vs RDF / Wikidata / Arango	Parsers, editors, SPARQL, dump tooling, sameAs graphs already exist for DBpedia. IEML has a demonstrator editor and papers.
Human authoring failed historically	Lévy now says LLMs should translate. Quality of <i>that</i> translation is unproven; garbage USLs would be a false sense of precision.
Third ontology risk	Personal graph + DBpedia + schema.org + IEML + Graph RAG entity titles = five names for one idea. Revival: do not freeze a giant schema. schema.org is the most respectable-looking freeze (§2.5).
Dictionary scale	~3000 words vs millions of Wikipedia topics. Coverage of <i>your</i> research (ecological redirection, ASC pivots, Drupal history) will be sparse unless you extend the dictionary — i.e. become an IEML ontologist.
IEML-Analysis-v07 complexity	The 2010s map (circuits, cités, bulb es , rhizomes) is exactly the “giant graph schema” the revival forbids as a first milestone.
Not ASC, not the UI codec	Putting morphemes in hashes or as ASC primitives contaminates the computational vocabulary with a linguistic theory. Already rejected.
Research software	Finite-state operations, rhizome algorithms, Intlekt editor: a second research stack beside Compose. Ops cost without users.
License friction	Lévy 2023 HAL is CC BY-NC-ND. The project’s code is GPLv2 / content CC BY 4.0. You can <i>read</i> the theory; you cannot casually fork that paper into a productized language runtime.
Does not solve KnowledgeGap	IEML can name a concept. It cannot tell you the concept is <i>unknown</i> , <i>insufficient for decision D</i> , or <i>out of budget</i> . That is the revival’s epistemic layer, missing from IEML-as-language.
Does not bind conditions to implementations	The Minimal Reasoning Model (requirements, fallbacks, environments) is a <i>task</i> calculus. IEML could annotate it; it does not replace <code>AND/OR/fallback</code> chains.
False precision	A USL looks more scientific than a Note. If the mapping is an LLM guess, you have laundered a probability into an identifier.

3.6 How to use it without implementing it (recommended)

Same hybrid as the 2026 memos: **some LLM remains the cognitive engine; which LLM is an ASC implementation; a symbolic layer is derived whenever the result must be handed to another provider or kept.**



Practical rules:

- 1. **Day one:** typed graph + uncertainty + KnowledgeGap + **which provider produced this step**. No IEML runtime.
- 2. **Pivot I/O is the handoff contract.** `run-agent` returns bounded artifacts, not “the model’s personality.” That is what makes Ollama, Tiiny, an API, and Cursor CLI interchangeable under one name.
- 3. **If a Concept keeps being paraphrased across years or across providers:** then consider a canonical id. That id may be a USL, a Wikidata QID (Q... item number), or an ASC-safe slug. Choose by *tooling*, not by loyalty. A QID is a *pointer into the public library*, not a reason to import Wikidata into Arango.
- 4. **Never** ask agents to “think in IEML”. Translate out, manipulate, translate back — if at all.
- 5. **If IEML is tried:** a Projet Complexe ASC pivot (`relate / annotate-concept`), same as code-graph-rag. UI still sees coordinates and titles.

Lévy vs this architecture, honestly:

Lévy 2023	Projet Complexe 2026
IEML as protocol of digital memory	ASC as protocol of <i>execution</i> ; Projet Complexe as <i>interpretation</i>
USL as URI	Entity ids + optional Wikipedia title / QID <i>pointers</i> already cover reference
Neurosymbolic integration	Several LLMs behind one pivot + graph + vectors, IEML optional
Collective-intelligence encyclopedia	Personal second brain + optional public grounding
Algebraic paradigms	Performance level of detail + genericity zoom (related <i>intuition</i> , different mechanism)

The overlap that is worth stealing: **semantic metadata as first-class, canonical form, do not let embeddings be the only memory.**

3.7 Verdict

IEML adds **tangible value only as a design compass and, later, as an optional annotation on durable Concept nodes**.
Multi-provider ASC does make explicit semantics urgent — not because Ollama “has no meaning”, but because Ollama’s meaning does not travel to Tiiny, to a remote API, or to Cursor CLI. The thing that travels is the typed graph (and provenance of who wrote it). That can be YAML and Arango. It does not have to be IEML.
It does **not** replace Graph RAG (corpus structure), DBpedia/Wikidata (reference), schema.org (a few public type *names*), Solr (words), or the task/knowledge killswitch (control).
Implementing the language before the typed personal graph — and before a pivot contract that names the LLM provider — would be a 2012-shaped mistake with 2026 compute bills.

4. LEGACY ENTITY DIAGRAMS → MERMAID


These diagrams (Histevents v04, Knowledge-oriented v08, Task-oriented v09, Minimal Reasoning Model v02) are the **pre-ASC, pre-LLM** data model. They still name the objects Graph RAG would extract, DBpedia would ground, schema.org would try to *type*, and IEML would try to coordinate.
Shared legend (all four):

Shape in draw.io	Meaning	Mermaid treatment below
Yellow ellipse	Node — content type	<code>(([node]))</code>
Blue rectangle	Term — taxonomy / vocabulary; often a <code>parent</code> loop	<code>[term]</code>
Pink ellipse	Relation — reified link with its own data	<code>{{relation}}</code>
Orange ellipse (task diagram only)	Custom entity — slice of an entity for implementation	<code>[[custom]]</code>
Arrow	Optional 1-1 or 1-n reference	<code>--></code>

They are **not** the 2026 storage architecture. They are the **interpretive vocabulary** Projet Complexe still owes the UI. Missing in all four, and required now: **KnowledgeGap / Unknown, confidence, provenance, agent, pivot**.
Originals: PNG + draw.io SVG in `assets/`.

4.1 Histevents — Entity Diagram v04

Historical-events domain: a **public-grounding-shaped** graph. This is what a *subset* of Wikipedia/DBpedia is for (`Event`, `Person`, `Place`, `Group`), not what your research journal is.



2026 reading: `Participation` is Graph RAG’s Relationship (needs its own data: role). `Thing (schema.org ?)` is **not** “import schema.org”. It is the *public kind* of Location/Group: optionally spelled as `schema:Person` / `schema:Place`, identified by a QID, looked up in the offline Wikipedia/DBpedia library — never loaded into Arango as the type system (§2.5). `Period.parent` is knowledge-mode *time zoom* (epoch → century → year) already described in the UI note. Do not LLM-extract Wikipedia to fill this; store a QID or a `frwiki/enwiki/ptwiki` title as a pointer.

4.2 Knowledge-oriented Entity Diagram v08


Personal epistemic model. This is the **second brain**. Graph RAG should *propose* `Links`; it must not own them.

Objective (original): bind **conditions** (and factors / priorities) to data for decision-making. Example: given requirements (environment, need a page callback yes/no, must/must-not), find the **closest match** among known implementations.

Simplification: **explicit chain of dependence** for requirements of implementations.


Open todo from the original: maybe skip a **Condition** entity and point Requirement straight at Technology — but then you lose single-responsibility for “elements of evaluation”.

Entity side:



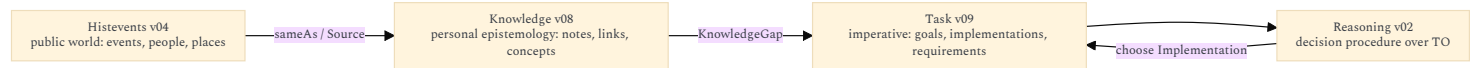
Note: requirements may be conditional: Implementation requires **X and (Y or Z)**.

Procedure the data model should generate:



When looking for fallbacks: mechanisms for discovering other alternatives — **tolerance, scope, relevance, value**. That tuple is the ancestor of Factor + the performance governor + the research budget on KnowledgeGap.

4.5 What the four diagrams are, together



2010s object	2026 mechanism
Source / Assembly / Note	Files + Solr + citations. Dublin Core names (title, creator, rights) are fine as field labels
Link + Type of Link + Factor	Schema-guided Graph RAG proposals, then human/agent-accepted Arango edges
Concept / Domain / Topic	Personal taxonomy; optional QID pointer; optional IEML USL later
Event / Person / Location / Thing	Offline Wikipedia/DBpedia/Wikidata library (fr/en/pt), not Arango nodes. schema.org at most spells the kind (Person), never supplies the instance
Task / Implementation / Requirement / Fallback	Agent control + killswitch; not an LLM latent graph
Comparison	Explicit pages (already a content type), not a cosine
Participation / Completion / Estimation	Reified relations with data — Graph RAG “covariates”

5. COMBINED STANCE (FOR LATER PIVOTS)



Confidence: the mapping to the existing architecture is high. Dump sizes and Graph RAG “winners” will move; the layering should not.

OPEN CHOICES (NOT DECIDED HERE)

- Kiwix ZIM vs raw XML vs DBpedia files vs a Wikidata truthy subset, for **offline** fr/en/pt lookup (never as an Arango import).
- Whether a personal node stores a QID, language-specific titles, or both, as pointers.
- Whether community reports are first-class Notes (knowledge-oriented) or generated sidecars Solr can search.
- Closed vocabulary for Graph RAG extraction: reuse Type of Link from v08, or a smaller revival set (supports, conflicts, unknown, sufficient-for). Public-world *entity* types: Histevents names, or a handful of schema.org URLs next to a QID — never the whole schema.org tree.
- Whether Requirement/Condition stay distinct (v02 todo) once agents generate fallback chains.
- If a Concept ever gets an IEML USL: store it as annotation, never as the only id.
- How `run-agent` names the provider capability (ollama / tiiny / api / cursor-cli) and whether Environment/Requirement from the task diagram gate lan-only vs remote.

Those can stay experiments behind `index` / `relate` / `research`. They do not change the rule:

ASC owns execution. Projet Complexe owns interpretation. Graphs, dumps, public ontologies, and IEML are mechanisms and coordinates — not a second control plane.