

OVERLAP WITH OPENRESEARCH

OPENRESEARCH VS ASC / PROJET COMPLEXE (EXPERIMENT TREE, NOT A SECOND BRAIN)

Date: 2026-09-09

Status: architecture note / design instrument (not a spec, not an implementation plan)

Reads: [alphaXiv/OpenResearch](#) (MIT, Rust, `orx` CLI; README, `SKILL.md`, `SYSTEM_PROMPT.md`, `agent-skills/orx-*`, `AGENTS.md`; v0.1.121 as of 2026-09-09), [openresearch.sh](#), sibling [openresearch-cli](#) (same product); ASC README (wrap, agent, builder, DSL, nested git, non-goals); Revival v4–v5 (NIH filter, paper mill, wrap Pi/Cursor); sibling overlap notes [EnvHarness](#), [AutoDesign](#), [Compile by Training](#); folder 1 review (*agentic-ai*, *rag*, *retrieval*, *packing*, *workflow*, *anonymization*) especially file 00 register D09, D14, D17, D22, D26–D29, D33–D35, D41 and file 06 (nilenso autoresearch, storymachine); folder 2 D06; Long *AI-Supervisor* (arXiv:2603.24402); Four layers; CLR note.

Repo created 7 Jun 2026. Desktop + CLI positioned as “the local-first workspace for research agents and autoresearch.” It wraps Claude Code, Codex, or OpenCode and installs a skill pack so those hosts can drive an **experiment tree**: literature, hypotheses, code variants, runs, logs, artifacts, optional LaTeX.

This note exists because the repo looks, at filename depth, like Projet Complexe’s research orientation — and it is not. It is the closest *shipped* cousin of a product v4/v5 already marked NIH: an autonomous “AI scientist” plus a paper mill. Steal the **experiment contract**. Do not become the workbench.

NAMES (DO NOT COLLAPSE)

Name	What it is	Relation here
OpenResearch (<code>orx</code>)	This product: local dashboard + CLI + skills + experiment tree + optional GPU marketplace	Subject of this note
openresearch-cli	Same codebase / same README; GitHub split	Same
nilenso autoresearch	Already on the shelf: <i>can an agent drive this CLI?</i>	Keep the word. OpenResearch’s README uses “Autoresearch” for a different loop (idea → code → run → next). Call <i>that</i> the experiment tree .
AlphaResearch (alphaxiv 2511.08522)	A paper: LLM algorithm discovery with a dual verifiable / simulated-review reward	Different system. Do not cite as this repo.

VERDICT

Strong overlap on **one axis** (a named, agent-drivable research CLI with frozen experiment nodes and evidence-in-logs), useful on a second (skills as interchange; wrap a coding agent), **conflict on product identity and on the LLM-economics hope**. OpenResearch puts a frontier coding agent *on* the whole research loop. Compile by Training and the ASC builder hope try to get LLMs *off* the hot path. Steal the four cardinal rules, “status is not evidence,” the stacked-bush tree, and the repair-then-ask cap. Do not import the workbench, the paper mill, parallel scientist-agents, the GPU marketplace, or SQLite as a knowledge store.

Dimension	Strength	Overlap	OpenResearch	Projet Complexe / ASC
Named, agent-drivable CLI	High	Strong	<code>orx projects / exp run / logs / paper / discover</code>	ASC pivots, DSL, <code>make \$subject-\$action</code>
Frozen node + fixed run command	High	Strong as policy	Four cardinal rules in <code>SKILL.md</code>	D22 host executes; comparable evals; git as history
Evidence in logs, not chat	High	Strong	<code>stdout</code> + <code>orx logs</code> ; status is not evidence	D33 traces; Winteringham; folder 2 event files
Wrap a coding agent, don’t be one	Med	Partial	Wraps Claude/Codex/OpenCode <i>and</i> becomes the dashboard	Wrap Pi/Cursor behind <code>run-agent</code>
Skills as files	High	High as interchange	<code>agent-skills/orx-*</code> installed into the host	YAML <code>able</code> → optional <code>SKILL.md</code> ; do not become the skill runner
Literature as retrieval	Med	Partial	<code>orx discover / orx paper</code> (alphaXiv, OpenAlex, bioRxiv)	<code>extract / research</code> ; hits are not Claims

Dimension	Strength	Overlap	OpenResearch	Projet Complexe / ASC
Local-first	Med	Split	SQLite + 127.0.0.1:4791; account for orgs/GPUs	Files SoR; Postgres for Claims; closable engines
Parallel agents / worktrees	Med	Conflict if copied as default	orx agent spawn; one helper, no nesting	Second agent only when a Requirement forbids the first
Reduce LLM use (builder / T0)	—	Opposite	Frontier agent on every round	Builder → T0; T3 rare
Paper mill / auto-scientist	—	Refuse as identity	orx-paper writes .tex; loop can run unattended	HITL commits Claims; publish is a contract
Modest hardware	Low	None	H100 catalog; 7-day runs	1050 4 GB; dedi is not an inference box
Knowledge ontology	Low	None	Experiment nodes, runs, artifacts	Claims, Links, Gaps, fr/en/pt

The overlapping sentence this repo shares with EnvHarness, AutoDesign, Compile by Training, and you:

Capability that survives a model swap lives in a named surround — not in a chat log.

Their surround is orx + skills + git branches + a coding-agent host. Yours is ASC names + PCA packs. Copy the CLI-as-vocabulary move. Do not copy the scientist in the seat.

WHAT OPENRESEARCH ACTUALLY IS

MIT License (copyright 2026 alphaXiv). Language: Rust (src/) + a committed ui/ dashboard. Homepage: openresearch.sh. Install:

```
curl -LsSf https://openresearch.sh/install.sh | sh
orx up
```


orx up opens a local dashboard at http://127.0.0.1:4791. Creating a project or launching a run does not publish code. An openresearch.sh account is for service-owned capabilities: organizations and **managed compute**. Official release builds send opt-out, coarse usage events tied to a random installation ID (no code, prompts, paths, repo names, tokens, emails, or experiment ids). Source/dev builds do not.

orx telemetry off exists.

AGENTS.md splits the product: orx owns the local CLI, dashboard, API, SQLite, coding-agent integrations, experiment orchestration, execution backends. openresearch.sh owns the website, docs, accounts, sandbox provisioning, managed-compute catalogs. Research projects, experiments, runs, logs, and artifacts are supposed to remain local.

Hosts. Claude Code (--append-system-prompt-file), Codex (developerInstructions), OpenCode (instructions list). SYSTEM_PROMPT.md is injected into every session: identity, project facts, artifacts path, evidence-and-links contract, skill routing. Task procedures live in native skills copied into the session worktree from agent-skills/.

Compute. Same committed snapshot can run locally, over SSH, or on Slurm, Kubernetes, Ray, Hugging Face Jobs, Modal, Tinker, and managed OpenResearch GPUs. Publishing the repo is not required for compute. orx up --remote user@host runs the workspace next to remote GPUs; the remote service binds to loopback **with no application-level authentication** — they document that other users on that host can reach it.



Skill pack (what the agent is taught)


Skill	Job	PC / ASC analogue
<code>openresearch-cli</code> (<code>SKILL.md</code>)	Cardinal rules + command index; <code>orx skill <name></code> loads modules	ASC entry-point index; genericity of docs vs core
<code>orx-experiment-tree</code>	Tree shape, freeze, repair cap, auto-research loop	Policy for variant-search Tasks — not a PC object
<code>orx-create</code>	<code>orx up</code> , <code>create-experiment</code> , baseline vs child	<code>make setup</code> ; entity create — different objects
<code>orx-git</code>	<code>orx/<slug></code> branches; commit before run; never rebase a frozen node	nested_git / worktrees; D06 git for harness files
<code>orx-evidence</code>	Design stdout; <code>orx logs</code> ; status $\neq$ evidence	D33 traces; eval plane (D41)
<code>orx-compute</code>	<code>orx exp run</code> only; never raw provider CLIs	ASC wrap: named pivot, not <code>ssh</code> in the model's mouth
<code>orx-instances</code>	Standalone instance in an org	remote_instance entity — refuse as control plane
<code>orx-lit-review</code>	Main agent ranks alphaXiv / OpenAlex / bioRxiv; do not delegate the retrieval loop	<code>research</code> retrieval; not Claims
<code>orx-paper</code>	Write real <code>.tex</code> in the working tree on first request	<code>publish</code> genre <i>preprint</i> — HITL, not auto
<code>orx-figures</code>	Publication figures / TikZ; default matplotlib is a defect	<code>publish</code> contract quality (Yu & Yao), not a plot library as identity
<code>orx-reports</code>	Artifacts directory: reports, CSVs, PDFs	<code>publish</code> drafts + <code>data/</code> — files as outputs
<code>orx-agent-delegation</code>	<code>orx agent spawn</code> ; no nested helpers; explicit compute authorization	Folder 1: second agent only when a Requirement forbids the first

The auto-research loop (as they specify it)

From `orx-experiment-tree`: toward a goal such as “best convergence for d=8”:

1. Read the baseline in the private worktree; find knobs in **code/config**, not in the command.
2. Form **one round**: co-equal options of a *single* decision.
3. Create that round as a **bush** under the right parent (baseline, or previous winner).
4. Edit each child's branch; commit; `orx exp run` with the **same** command.
5. Read logs (`orx-evidence`); pick a winner; descend.

Do not edit a frozen node. Do not rewrite the run command.



Wrong shapes they name: **flat fan** (every variant off the root — wins never accumulate) and **noodle** (a long single-child chain of co-equal options). Right shape: **stacked bushes**.

Literature loop

`orx-lit-review`: the *main* agent is the retrieval ranker. Primitives (no login): `orx discover keyword|embedding|openalex|biorxiv`, then `orx paper <id>`. Difficulty 1–10 sets follow-up budget (0 / 1 / 2 rounds). Dedup by id, DOI, arXiv id, then title. Prefer alphaXiv representation of an arXiv duplicate because it supports full-text. Scholarly claims use those source links, not project file tags. This is a real retrieval skill. It is still **retrieval**.

Paper loop

`orx-paper`: do not answer a paper request with an outline in chat. Create `paper.tex` at repo root (or `<topic>.tex`). Prefer `.orx/latex-templates/` if present. Figures go through `orx-figures`. LinkedIn launch copy (Aug 2026) sold “idea → paper end-to-end” and `/write-paper`. That is AutoDesign’s mill with a different medium (LaTeX vs poster HTML).

FOUR CARDINAL RULES (THE STEAL)

From the top-level `SKILL.md`. They say breaking any one **silently invalidates results** — not style preferences.

1. **Never edit a node once a run has answered it.** Freeze is permanent; a disappointing number is still a result. Until a run answers, the node is **provisional** (repair in place). New idea → **child**.
2. **The run command and the environment are a fixed contract** — identical on every node. Children inherit the parent’s command verbatim. No `LR=3e-4 python ...`. Set once: `orx project edit <id> --run-command '...'`.
3. **Vary committed code/config, not knobs in the command.** Same command, different git, so logged summaries stay comparable.
4. **Grow the tree downward, not sideways.** Fan a little *within* a round, then descend onto that round’s winner. A root with a long row of direct children and no grandchildren is the failure mode.

Thin killswitch alongside:

- **Repair cap**: two runs in a row that answer nothing on one node, then **ask the user**. Same failure on a second node → ask (setup problem). Bare relaunch or flavor/backend switch counts as a repair.
- **Scientific stop**: “~3 failed or regressed runs” (mentioned as separate from repair).

Evidence rule (`orx-evidence`): print metrics and config on stdout; before reporting a run-derived claim, read the log; **run status is not evidence**; truncated output is not evidence of absence.


Delegation rule (`orx-agent-delegation`): helper gets its own worktree and transcript; cannot spawn another helper; never give it a branch this session has checked out; **state exactly which** `orx exp run` **calls it may launch**; never delegate the literature ranking loop.



FIVE-WAY: WHAT EACH SIBLING FREEZES

Do not treat OpenResearch as a fourth *compiler*. It is a **runtime host for experimental coding**. The August/September papers freeze a core and rewrite a surround. OpenResearch freezes *answered experiment nodes* and keeps rewriting *children* with a live frontier agent.

	EnvHarness	AutoDesign	Compile by Training	OpenResearch	Projet Complexe
Frozen core	env + verifier	model θ	0.6B interpreter + spec	answered node's commit	ASC names + human goal / Claim
What runs on the hot path	policy in a gym	designer-critic	<code>Run(p_s, x)</code> on 0.6B	Claude/Codex/OpenCode	named pivot; T0 if possible
Outer loop	EnvRigger writes wraps	P patches H	teachers + LoRA at compile	auto-research tree	optional gated pack edits, HITL
Success metric	held-out SR	PosterBench	FuzzyBench LEM	EVAL / logs / a <code>.tex</code>	ingest, curate, export, research; fr/en/pt
SoR	Python subclasses	harness source + Skills	<code>.paw</code> + spec	git branch + SQLite + artifacts	YAML + DSL; files for bytes
LLM economics	designer spends models	outer P spends models	teachers at compile only	models on every round	builder → T0; T3 rare




EnvHarness answers: *reshape what the agent sees*.
AutoDesign answers: *reshape the software that produces the next artifact*.
Compile by Training answers: *compile a recurring function and stop calling the teacher*.
OpenResearch answers: *give a coding agent a comparable experiment tree and let it climb*.

Projet Complexe needs OpenResearch’s **comparability rules** as policy for variant-search Tasks, and must not take the climb as the research orientation.

CONCEPT MAP (NOTES → OPENRESEARCH)

Idea in your notes	OpenResearch counterpart	Overlap
ASC wraps frozen CLIs; names persist	<code>orx</code> is the named CLI; hosts swap (Claude / Codex / OpenCode)	High as a pattern. They still <i>are</i> a dashboard. You wrap, you do not become.
v5 NIH: do not rebuild Pi/Cursor; wrap	They wrap those hosts <i>and</i> add <code>orx up</code>	Partial. Steal wrap. Refuse second IDE.
v5 NIH: paper mill	<code>orx-paper</code> + launch copy “idea → paper”	Conflict as identity. <code>publish</code> genre only.
Long AI-Supervisor: steal gaps, refuse auto-commit	Experiment tree proposes next nodes; no Claim object	High as a warning. Same refuse.
nilenso autoresearch: drivability of a CLI	<code>orx</code> is an agent-facing CLI with skills	High. Use as a <i>specimen</i> of drivability, not as the vocabulary. Keep the word “autoresearch” for nilenso.
storymachine: spec as object	experiment <code>--description</code> is the hypothesis	Low–medium. Description ≠ acceptance criteria on a product Task.
D22 host executes; model proposes	<code>orx exp run</code> only; never raw <code>python train.py</code> as an untracked job	High. Same wrap discipline.
D24 DSL / do not teach frontier a private dialect	skills teach <code>orx ...</code> as the dialect	Medium. Better than a private shell; still a second vocabulary beside ASC. One Implementation behind a pivot, if ever.
D09 killswitch / System M	repair cap; scientific stop; spawn compute must be authorized	Medium as budget. Not orientation (research may not <code>bash</code> the household).
D27 lethal trifecta	lit fetch + private repo + <code>exp run</code> in one agent catalog	Conflict unless catalogs split. Spawn’s explicit launch allowlist is a start.
D28 HITL on Claims and harness patches	ask user after two failed repairs; paper writes itself	Low. Repair HITL ≠ knowledge HITL.
D29 local-first; closable	local SQLite by default; remote loopback has no auth; GPU marketplace optional	Split. Steal local default. Refuse remote-without-auth and marketplace-as-identity.
D33 traces	<code>orx logs</code> ; file/run tags in chat	High as discipline. Different envelope than folder 2 traces.
D06 files as SoR	git branch is SoR for experiment <i>code</i> ; SQLite is live index	Medium. SQLite as derived experiment index is fine. Not for Claims.
D17 cascade; modest hardware	H100 catalog; 7-day max run	None as environment. A <code>run-command</code> that is <code>make test</code> could run on the laptop; their demos are nanochat / GRPO.
D26 MCP is transport	CLI-first, not MCP-as-brain	Confirmed (accidentally).
D35 refuse-list / AI-scientist	the product	Refuse.
Folder 1: most “multi-agent” is a workflow; second agent only when a Requirement forbids the first	<code>orx agent spawn</code> : one helper, no nesting, standalone brief	High as discipline if kept rare. Conflict if parallel scientists are the default.
Compile by Training: fewer runtime LLM calls	opposite	Use as contrast , not as merge.

Idea in your notes	OpenResearch counterpart	Overlap
AutoDesign: evolve H , gated	they evolve <i>code under experiment nodes</i> , not the <code>orx</code> harness, autonomously	Conflict on gate. AutoDesign’s train/dev + HITL on H is stricter than their tree climb.
CLR / Flow	worktrees isolate context; they do not measure a band	Low. Isolation $\neq$ packing.
Four layers	execution + a dynamics loop on the tree	Low as ontology. No Claim.
fr/en/pt as Requirements	English LaTeX / arXiv-first discovery	None.
Builder / T0 hope	skills so the LLM can drive <i>more</i>	Opposite.



REGISTER CROSSWALK (FOLDER 1 D-IDS)

ID	Decision	This repo	Verdict
D01	Three scopes; coupling is the product	One workbench + coding-agent host	Silent on PC meaning; conflict if ASC becomes <code>orx</code>
D06	Files SoR; derived rebuildable	Git branch SoR for experiment code; SQLite live	Steal the split for <i>experiments</i> ; keep Claims in Postgres + file export
D09	Orientations + killswitch	Repair/scientific stops; no knowledge/task split	Silent on orientation; do not let <code>research</code> inherit <code>exp run</code>
D14	Harness not weights	Skill files + playbook are the harness; they do not freeze the <i>coding</i> model	Confirmed as surround ; the surround is their product, not yours
D17	Cascade tiny → LAN → remote	Frontier host by default; GPU catalog	Contradicted as default ; a local <code>run-command</code> could still be T0/T1
D22	Host executes	<code>orx exp run</code> is the only legal launch	Confirmed — cite the wrap
D24	DSL compiles to schema	Skills teach <code>orx</code> argv	Adapt as one Implementation, not a second primordial vocabulary
D26	MCP is not the control plane	CLI + local API; dashboard is a local UI	Partial. Local UI is closer to Tauri-thin than to messenger-as-host. Still: ASC/PCA over Compose remains the household control plane.

ID	Decision	This repo	Verdict
D27	Lethal trifecta	Lit + private repo + outbound compute in one catalog	Conflict unless split
D28	HITL for knowledge and harness	Repair ask; paper without accept	Not a substitute
D29	Local-first; closable	Local default yes; remote no-auth; marketplace	Split
D33	Traces first-class	Logs + tags	Confirmed as discipline
D34	Redirection drills	how_to_turn_off for a Technology orx would be: stop wrapping; pivots remain	Yes if orx is never identity
D35	Refuse-list	AI scientist / mill	Refuse the product
D41	Eval plane	EVAL.md / stdout metrics (docs + evidence skill)	Steal the artifact shape (a run must write comparable evidence)

Folder 2: an experiment branch is a *fonds* of a trial (code + commit + log). That is Edwards / Servais-adjacent and good. It is not a Claim. Promotion to L3 still needs HITL.

VERSUS THE HOPED-FOR ROADMAP (BUILDER → T0)

Hoped-for move	OpenResearch
Builder fills tooling gaps so fewer LLM calls	Skills so the LLM can drive more of the loop
Map NL → named entry points	They do this <i>for orx</i> , then keep a frontier model in the seat
DSL + rules as combinators	Cardinal rules combinators for experiments , not household Tasks
Fewer human validations	Fewer babysitting steps on <i>runs</i> ; no human on meaning
Bigger models rare	Bigger models are the product

If OpenResearch were installed as the research orientation, mass would move **away** from T0, not toward it. Compile by Training is the ally of that hope. This repo is the counterexample.

STEAL

Cite OpenResearch as a **shipped experiment workbench** that operationalizes comparability: freeze answered nodes, fix the run command, vary git, shape the tree as stacked bushes.

Steal the **four cardinal rules** as PCA policy for any Task that is “try variants on a repo” (harness-pack eval, packing ablation, client A/B). You do not need *orx* to keep a Makefile target identical across branches.

Steal **status is not evidence**. Map onto D33 and onto autoresearch class C (silent wrong / incomplete). A pivot that exits 0 without writing the metric has failed the contract.

Steal **repair cap then ask**. Two unanswered runs, then drama (Lefèvre), not a third silent retry. Separate “setup is broken” from “the hypothesis lost.”

Steal *orx exp run* **as wrap**: the model must not invoke the trainer/provider CLI. ASC already wants this (named *\$action*). Their compute skill is a worked example of that sentence.

Steal **do not delegate the retrieval ranking loop** to a sub-agent. Main agent inspects candidates. That matches folder 1’s “coordination is paid in context” and the killswitch against swarms.

Steal **skills as interchange** (again): they install *SKILL.md* into Claude/Codex/OpenCode. Same AutoDesign move. Emit from YAML *able*; do not make ASC a skill runner.

Steal **git worktree per session** as an *Implementation* for isolation of code Tasks — when the user asked for a worktree, and never as unsolicited gitflow. Frozen branch = immutable archive of what ran.

Steal **EVAL.md / compact stdout summary** as the eval-plane artifact a scientific or harness-eval run must write (D41). Winteringham: the checker reads that file, not the chat.

Steal spawn’s **explicit compute authorization** in the helper brief. Catalog-level trifecta needs more; this is the minimum for a second session.

REFUSE

Do not make OpenResearch the research orientation, the Tauri replacement, or the ASC dashboard. Wrap a harness; do not become a scientist IDE.

Do not take “idea → paper end-to-end” as `publish`. A `.tex` is one genre Implementation under HITL and a quality contract. Nolan: revision is the craft. AutoDesign overlap already refused the mill.

Do not run the auto-research loop unattended on household knowledge, client repos, or ASC core. Legal objects for a *gated* climb, if any, are **packs / eval recipes / experiment code the human named** — not Claims, not primordial YAML, not the user’s Task.

Do not collapse “Autoresearch” with nilenso. Different loops. Folder 1’s word stays: drivability of a CLI.

Do not put `orx` in ASC core. Optional later: a contrib Technology behind `exp-run` / `exp-logs` if a named Task is blocked without an experiment tree (v5: wrap after two foreign tools have needed it). Until then, NIH / tool obsession (Osmani).

Do not share one catalog among `orx discover` (untrusted web), a private corpus, and `orx exp run` (outbound compute). D27.

Do not use `orx up --remote` on dedi-2025: loopback, no application auth, other users on the host. That fails D29 and the dedi threat model.

Do not use managed OpenResearch GPUs as identity. H100 marketplace is not the 1050 / Tiiny / closable-engine story. Overflow compute, if ever, stays a named Technology with quota, redaction, and `how_to_turn_off`.

Do not treat SQLite as the second brain. Derived index for experiments only.

Do not train or fine-tune because a demo is nanochat / Qwen GRPO. D14 / D42 unchanged.

Do not let parallel `orx agent spawn` become multi-agent consensus as truth (2603.24402). Helper writes a brief’s deliverable; human commits meaning.

Do not send the household literature graph through alphaXiv as SoR. `orx paper` is a retrieval Implementation; extract-once and Claims stay yours. English-first discovery does not satisfy D07.

CAVEATS (SO YOU DO NOT OVER-CITE)

Issue	Why it matters here
README “local-first” vs marketplace + <code>orx login</code>	Local <i>default</i> is real; attachment to <code>openresearch.sh</code> exists for the features they will demo.
Telemetry on official builds	Opt-out, coarse; still a channel. Source builds off.
Remote no-auth	Documented; easy to miss.
“Autoresearch” naming	Collides with nilenso.
Paper skill vs launch marketing	Skill is “write <code>.tex</code> in tree”; marketing is end-to-end automation. Judge the skill; refuse the slogan.
No household eval	Nanochat / GRPO templates are not ingest-claim-publish in fr/en/pt.
Skills assume a frontier coding agent	Drivability of <code>orx</code> is not drivability of ASC on a 1.7B.

WHERE THE ANALOGY BREAKS

OpenResearch is a **training-and-trial workbench for people who already live in a coding agent and a GPU queue**. Projet Complexe is a **runtime cognitive institution**: human-judged knowledge plus a named computational vocabulary. ASC is **glue**.

That changes “research.” In their setting, research is stacking experiments on a codebase until a metric moves, then writing LaTeX. In yours, research is extract → propose → HITL → Claim / Gap, with a killswitch against acting. An experiment tree can be a *Task type* (code/science Implementation). It cannot be the knowledge orientation.

That also changes “local.” Their local is “SQLite on loopback, code not published.” Yours is files as SoR, engines closable, Three Safes on any hop, dedi as GitOps not as an open `orx` bind.


Four-layer note: almost entirely **execution**, with a **dynamics** loop on the tree. No ontology of household objects. Cognitive institutions: they improve the *procedure that produces experiment artifacts*; Superpowers-like skills are what they ship; Guardrails are thin (repair cap, spawn allowlist). Complementary if kept in the **code-experiment worker** slot. None of Projet Complexe’s knowledge institution.

Flow / CLR: a worktree lowers *interference* between sessions. It does not pack a corpus or keep a 7B in band. EnvHarness still owns the band metaphor. Compile by Training owns amortized cost. This repo owns **comparable trials**.

PLACEMENT ON THE THREE-PROJECT CUT

Project	Relation to OpenResearch
ASC	Cousin of wrap-don’t-rebuild. Document <code>orx</code> as a Technology a pivot might call, like Cursor CLI. Do not add an experiment-tree runtime or a paper compiler to core. Optional later: generate <code>SKILL.md</code> that teaches ASC pivots the way they teach <code>orx</code> .

Project	Relation to OpenResearch
Projet Complexe ASC	If anything is stolen, it is policy : frozen answered trials, fixed run command, stacked bushes, status ≠ evidence, repair-then-HITL, retrieval ranking not delegated, spawn with an explicit compute allowlist. <code>EVAL.md</code> as an eval-plane contract. <code>research</code> stays budgeted and catalog-split.
Projet Complexe	The Flow UI remains the interpreter of Tasks and Claims. A literature hit is retrieval. A run metric is evidence for a <i>proposed</i> Claim. A <code>.tex</code> is a <code>publish</code> genre. No new PC object named Experiment — or, if one appears later, it is a Task subtype, not a Claim.



Genericity test (v5 §4.4): would another project, not a second brain, need this? **Yes** — any repo that runs comparable trials (ML, compiler flags, harness evals). That is an argument for keeping the *rules* in PCA policy. It is not an argument for `orx` as the household shell. Another project needs `make test-integration` to mean the same thing on every branch, not `orx exp run` as the name of work.

PRACTICAL TAKEAWAY

Read OpenResearch next to the three overlap notes, not instead of them.

1. **EnvHarness** — freeze the world, wrap `reset/step`, regulate difficulty.
2. **AutoDesign** — freeze the model, evolve *H*, gate with a frozen eval.
3. **Compile by Training** — freeze a tiny interpreter, compile NL, spend teachers on the build.
4. **OpenResearch** — freeze *answered trials*, wrap a coding agent, climb a git tree, optionally write LaTeX.
5. **Projet Complexe** — freeze names and the human’s objective; compile gaps into **files**; comparable trials as *policy* when the Task is variant-search; only a human commits Claims.

The sentence to keep from this repo:

Never edit a node once a run has answered it. Vary code, not the command.

Your node is a git-recorded Implementation of a named pivot. Theirs is an `orx` experiment. Same sentence. Different catalog.

Sources: OpenResearch README (local-first, hosts, compute, telemetry, `orx` commands); `SKILL.md` (cardinal rules); `SYSTEM_PROMPT.md` (playbook, evidence tags); `AGENTS.md` (`orx` vs `openresearch.sh`); `agent-skills/orx-experiment-tree`, `orx-create`, `orx-git`, `orx-evidence`, `orx-compute`, `orx-lit-review`, `orx-paper`, `orx-reports`, `orx-agent-delegation`, `orx-figures` (via GitHub raw, 2026-09-09); GitHub metadata (created 2026-06-07, MIT, Rust, v0.1.121); ASC README; Revival v4–v5 NIH / wrap / paper mill; folder 1 file 00 §§12, 16.2, 16.4, 16.7 and file 06 (autoresearch, storymachine); folder 2 D06; Long 2603.24402; sibling overlap notes.

- github.com/alphaXiv/OpenResearch
- github.com/alphaXiv/openresearch-cli
- openresearch.sh
- Siblings: [Overlap with EnvHarness](#), [Overlap with AutoDesign](#), [Overlap with Compile by Training](#)
- [project-complexe](#)

- [projet-complexe-asc](#)
- [asc](#)