

PI, OMP, SUPERPOWERS, SKILL PACKS

Date: 2026-09-09

Status: architecture note / design instrument (not a spec, not an implementation plan)

Checked: GitHub APIs and upstream READMEs / docs on 2026-09-09. Star counts move; treat them as *attention*, not quality.

Reads: [earendil-works/pi](#) (MIT, TypeScript; README, packages/coding-agent/docs/{skills,packages,extensions,rpc, sdk, containerization}.md, [pi.dev/docs/latest](#)); [can1357/oh-my-pi](#) (MIT, TypeScript + ~80k LoC Rust natives; README, [omp.sh](#), docs on skills / custom tools); [obra/superpowers](#) (MIT; README, skills/*/SKILL.md, Pi install path); Agent Skills standard ([agentskills.io](#)); sibling notes in this folder; Revival v4-v5; folder 1 (*agentic-ai, rag, retrieval, packing, workflow, anonymization*) especially files 00, 01, 02, 06, 07, 09, 10; folder 2 (*data stores...*) files 01, 04, 06.

This document **paraphrases**. It does not paste books, skill bodies, or chat dumps into the second brain.

Each block is asked the same four questions used on this shelf:

1. What was actually claimed (repo, spec, or marketplace)?
2. How is it implemented in the wild?
3. Where does it sit on ASC / Projet Complexe ASC / Projet Complexe / folder 1 / folder 2?
4. Steal, adapt, or refuse?

WHY THIS NOTE EXISTS

v4 already named Pi as the coding-agent **worker** not to rebuild. Folder 1 already named `SKILL.md` as an **export** of YAML `able`, not a store. Folder 2 already put `harness/SKILL.md` next to `able.yml` in the file SoR. Superpowers is already *running in this Cursor workspace* as a process overlay.

What the three GitHub URLs plus the 2026 skill-pack market add, that those notes had not put in one instrument:


1. **Pi moved from a small Mario Zechner repo into Earendil Inc.** (`@earendil-works/*`, `pi.dev`, packages, RPC, Agent Skills as a first-class loader). The wrap target is still the same object. The distribution surface is larger.
2. **Oh My Pi is not “Pi with a nicer TUI.”** It is a batteries-included *fork* that adds IDE-grade tools, a memory bank the agent writes, `learn / manage_skill`, browser / desktop / web search, ACP, and a 31-tool catalog. That is a different Meadows spend: rules and goals, not parameters.
3. **Superpowers is not a harness.** It is a *methodology encoded as mandatory skills*. It can ride Pi, OMP, Cursor, Claude Code, Hermes, ... without owning execution.
4. **Skill packs became a distribution layer** (Anthropic plugins, Vercel `skills.sh` packs on 2026-08-07, awesome lists, vendor marketplaces). That is an interchange *and* a supply-chain attack surface. It is not a knowledge ontology.

The evaluation question is not “which of these should Projet Complexe become?” It is: **which of these can sit behind named pivots without becoming identity, and which attachments they would close.**

NAMES (DO NOT COLLAPSE)

Name	What it is	Relation here
Pi (<code>earendil-works/pi</code> , formerly <code>badlogic/pi-mono</code>)	Minimal terminal agent harness : unified LLM API (<code>pi-ai</code>), loop (<code>pi-agent-core</code>), coding CLI (<code>pi-coding-agent</code>), TUI, telemetry. MIT. ~103k GitHub stars (2026-09-09).	Wrap target. Technology behind <code>run-agent</code> for <i>code</i> Tasks. Not the second brain. Not the control plane.
Oh My Pi / OMP (<code>can1357/oh-my-pi</code>)	Fork of Pi by Can Bölük / Stencil Labs. Same package <i>shape</i> , rewritten coding-first: LSP, DAP, hashline edits, 31 tools, memory bank, subagents, ACP, collab relay. ~30k stars. Site: omp.sh .	Alternative / later Implementation of the same <code>run-agent</code> worker. Stronger capability, stronger identity trap.
Superpowers (<code>obra/superpowers</code>)	Jesse Vincent / Prime Radiant. Process skill pack: brainstorm → plan → TDD → review → finish. MIT. ~284k stars. Installs into 14+ harnesses including Pi (<code>pi install git:github.com/obra/superpowers</code>).	Procedure overlay. Closest cousin of PCA Requirements + HITL gates. Not tools. Not Claims.
Agent Skills (<code>SKILL.md</code>)	Open standard (Anthropic origin; agentskills.io): folder + YAML frontmatter <code>name / description</code> , progressive disclosure, optional scripts/references.	Interchange format. Folder 1 D24: generate from YAML <code>able</code> ; do not become the skill runner.
Skill pack	A <i>bundle</i> of skills (git repo, Claude plugin, Vercel pack URL, Pi package).	Distribution unit. Not a pivot. Not a Claim schema.

Name	What it is	Relation here
Pi package	npm/git bundle of extensions + skills + prompts + themes (<code>pi install ...</code>).	How Superpowers rides Pi. Analog of a PCA composition, not of a PC object.
Hermes Agent (<code>NousResearch/hermes-agent</code>)	Self-hosted harness with the packing tricks v5 already stole. Loads Superpowers as a plugin. ~244k stars.	Packing reference. Refuse gateway-as-host. Already treated in About Memory, RAG, and Graphs .
ECC (<code>affaan-m/ECC</code> , formerly “everything-claude-code”)	Full overlay: skills, instincts, memory, security, research-first. MIT + commercial GitHub App. GitHub API ~255k stars (treat as attention, possibly inflated).	Competitor <i>bundle</i> to Superpowers + memory products. Do not import as identity.
OpenResearch (<code>orx</code>)	Research CLI + skill pack that <i>wraps</i> a coding host.	Already a sibling note. Skills-as-interchange cousin. Refuse paper mill.



A skill pack never sits above ASC. If it starts writing Claims, memory banks, or catalogs, it has inverted the stack.

0. VERDICT IN ONE PAGE

Used (workers + interchange), refused (product identity). Pi is still the right *shape* of coding worker: small core, typed tools, lifecycle events that ASC can mirror, Agent Skills loader, RPC/SDK so a host can start it. OMP is the same shape with the IDE and a private memory OS bolted on — steal headline / LSP *as optional tool Implementations*, refuse the memory bank, `learn`, `computer`, default `web_search`, and collab relay as household infrastructure. Superpowers is the right *shape* of process brief: mandatory HITL before implementation, plans as objects, verification before “done.” Steal the *gates*; do not let Superpowers own git branches, worktrees, or the research killswitch. Other packs are a catalog to *read*, not a dependency to *become*. The overlapping sentence this shelf shares with EnvHarness, AutoDesign, OpenResearch, and v5:

Capability that survives a model swap lives in a named surround — not in a chat log.

Their surround is a coding harness + `SKILL.md` directories + (for OMP) a SQLite memory engine. Yours is ASC names + PCA packs + files/Postgres. Copy the **event list, progressive disclosure, and process gates**. Do not copy the scientist, the memory OS, or the marketplace.

Dimension	Strength	Overlap	Upstream	Projet Complexe / ASC
Named loop around frozen weights	High	Strong	Pi <code>transformContext</code> → <code>convertToLlm</code> ; OMP same plus stream-abort rules	<code>pre_llm</code> / <code>post_llm</code> ; packing is code (D16)
Tools as schemas the host executes	High	Strong	Pi <code>registerTool</code> ; OMP 31 named tools	YAML <code>able</code> → JSON Schema; host executes (D22)
Skills as on-demand briefs	High	High as interchange	Agent Skills; Superpowers; Pi/OMP loaders	Generate <code>SKILL.md</code> ; do not become the runner (D24, D50)
Progressive disclosure	High	Strong	descriptions in system prompt; body via <code>read</code>	D13; Hermes <code>skills_list</code> → <code>skill_view</code> ; LOD
Lifecycle hooks	High	Strong as names	Pi/OMP events; Claude Code exit codes (folder 1 §29)	Filename-addressable hooks; fail-closed on gates
Multi-provider + local	High	Adapt	Pi <code>llama.cpp</code> ; OMP Ollama / custom <code>models.yml</code> (Tiiny-shaped)	Technology YAML; laptop → Tiiny → overflow

Dimension	Strength	Overlap	Upstream	Projet Complexe / ASC
Process before code	High	Adapt as HITL	Superpowers brainstorm / plan / TDD / review	HITL commit; spec pivot (D38); killswitch
Agent-written memory	High (OMP, Hermes, ECC)	Conflict	retain / learn / MEMORY.md	Claims via HITL; files + brain/events SoR (D06 v6)
Fat tool catalog	High (OMP)	Conflict if default	31 tools incl. browser, computer, web_search	Catalog is a design artifact; trifecta (D27)
Subagent fan-out	Med	Narrow	OMP task; Superpowers SDD	Second context when window or Requirement forbids the first; not multi-agent ontology (D35)
Marketplace / packs	Med	Refuse as SoT	plugins, npx skills add, Pi packages	YAML able is SoT; packs are optional export/install
Modest hardware / CLR	Low-Med	Split	Frontier APIs default; local is a provider among sixty	1050 4 GB; packing + triage; tiny local first
Knowledge ontology	None	None	files/diffs/sessions	Claims, Links, Gaps, fr/en/pt
Control plane	—	Refuse inversion	TUI / ACP / collab / messenger	ASC starts engines; Tauri is thin; no daemon-as-host

Recommendation for the evaluation stage (not a build order):

1. Keep **Pi (thin)** as the default *candidate* coding worker: RPC/SDK, --no-skills possible, no built-in memory OS, events you can map. Stage 1 may still ship *only* Ollama + Cursor CLI + a 50-line loop (v5 \$9 still open); Pi is the thing you wrap *if* you wrap a coding harness.
2. Treat **OMP as a later, allowlisted Implementation** of the same pivot — after a catalog recipe exists that can pin --tools read,edit,bash,... and turn off memory / browser / computer / github. Do not start from OMP’s default box.
3. Treat **Superpowers as a source of gates**, not as the household methodology. Translate brainstorming → design HITL; writing-plans → spec / Requirement pack; verification-before-completion → eval plane (D41); TDD → code-Task Requirement, not research-Task. Do not let it create worktrees or branches as a default (this repo’s git rules already forbid unsolicited branches).
4. **Do not install popular skill packs into the knowledge corpus.** Discover on skills.sh / awesome lists; copy a folder only after reading scripts; pin a hash; record provenance (folder 1 §27). YAML able remains the source of truth.

JARGON NOTES

Term	Alternative notations	Definition in this note	Examples
Harness	surround; AutoDesign H; agent loop	Everything around frozen weights: prompts, tools, validators, packing, killswitch	Pi; OMP; Cursor Agent; a PCA wrap
Worker	coding agent; Technology behind run-agent	Process ASC <i>may start</i> . Does not own Claims or the control plane	pi; omp; cursor CLI
Skill	Agent Skill; SKILL.md; procedure brief	On-demand instructions (+ optional scripts). Not a tool, not a Claim, not an OS	Superpowers brainstorming; a generated extract brief
Tool	function calling; action	Named operation with a schema, executed by the host after the model requests it	Pi read; OMP lsp; ASC entry point
Extension	Pi/OMP TypeScript module	Code that registers tools, commands, event handlers	ASC→registerTool adapter would be one
Pi package	pi install git:... / npm:...	Bundle of extensions + skills + prompts + themes	Superpowers on Pi
Progressive disclosure	LOD for skills	Names/descriptions always; bodies on demand	Pi XML skill list; Huang “router not repository”
Hashline	OMP edit format	Patch by content-hash anchors instead of exact string replace	Tool Implementation, not a knowledge format

Term	Alternative notations	Definition in this note	Examples
Memory bank (OMP)	retain / recall / Mnemopi	Agent-curated facts in SQLite (or Hindsight)	Not Claims. Episodic at best; refuse as SoR
Managed skill (OMP)	learn → manage_skill	Worker writes a skill under ~/.omp/agent/managed-skills	Conflicts with D50 (human-only skill writes)
ACP	Agent Client Protocol	Editor ↔ agent JSON-RPC (Zed et al.). Distinct from MCP	omp acp; folder 1 already parked this as a Technology option
Lethal trifecta	Bhagwat	Private corpus + untrusted web + outbound tools in one catalog	OMP default web_search + bash + your PDF tree
Trigger testing	Montaldo ch. 4	Did the model load the right skill at all?	Eval for description quality; not “did the code compile”
Skill pack (Vercel)	npx skills add pack URL	Unlisted bundle of several skills, not ACL-protected	Team sync mechanism; supply chain = URL leak

If a vendor word is missing here, treat that as a hint: it probably should not become identity.

1. PI (earendil-works/pi)

1.1 WHAT IT ACTUALLY IS

MIT. Language: TypeScript (Node/Bun). Created 2025-08-09. Homepage/docs: [pi.dev](#). npm: [@earendil-works/pi-coding-agent](#). GitHub: ~103,468 stars, ~12,923 forks (2026-09-09).
The repo README’s one-line job: **AI agent toolkit** — unified LLM API, agent loop, TUI, coding agent CLI. Docs call it a **minimal terminal coding harness**, small at the core, extended through TypeScript extensions, skills, prompt templates, themes, and Pi packages. [badlogic/pi-mono](#) now resolves to this repo. Oh My Pi still cites Mario Zechner’s Pi as the parent. Earendil Inc. is the current publisher ([earendil.com](#) on the docs footer). That is a *publisher* change, not a different architecture.

Packages (upstream table):

Package	Role
@earendil-works/pi-ai	Multi-provider LLM API (OpenAI, Anthropic, Google, ...)
@earendil-works/pi-agent-core	Loop, tool calling, state
@earendil-works/pi-coding-agent	Interactive CLI
@earendil-works/pi-tui	Differential terminal UI
@earendil-works/pi-telemetry	Vendor-neutral telemetry contracts
@earendil-works/chord	Application-composition runtime (services, replicated state, RPC, plugins)

Sibling: [earendil-works/pi-chat](#) for Slack/chat automation — **messenger-shaped**. Out of identity (v5 refuse messenger-as-host).
Install (documented):

```
npm install -g --ignore-scripts @earendil-works/pi-coding-agent
# or
curl -fsSL https://pi.dev/install.sh | sh
```

They treat npm lifecycle scripts as a supply-chain issue: exact pins, lockfile as ground truth, `npm ci --ignore-scripts`, shrinkwrap for the published CLI. That discipline is closer to folder 3’s intended SBOM-for-AI than most agent READMEs.

1.2 HOW IT IS IMPLEMENTED

Documented pipeline (already in Revival v4 §3.1):

```
AgentMessage[] → transformContext() → AgentMessage[] → convertToLlm() → Message[] → LLM
                  (optional)                                (required)
```

`transformContext`: prune, inject. `convertToLlm`: drop UI-only messages, fit the provider.

Lifecycle events the coding agent exposes (map, do not reimplement):

Pi event	ASC analogue
before_agent_start	pre_run-agent / pre_llm
before_provider_request	pre_llm variant per provider
after_provider_response	post_llm
tool_call (block or modify)	pre_action — authz lives here
session_before_compact	pre_compact
registerTool + schema	YAML able → JSON Schema

Permissions. Pi does **not** ship a permission sandbox. It runs as the launching user. Documented isolation patterns: Gondolin (tools in a micro-VM, auth on host), plain Docker, OpenShell. Folder 1 already stole “Docker the foreign tools” (D30) and inverted Claude Code’s fail-open on hook error for boundary gates.

Skills. Implements the Agent Skills standard, lenient on the “folder name must match `name`” rule because shared directories across harnesses make that rule painful. Discovery:

- Global: `~/.pi/agent/skills/`, `~/.agents/skills/`
- Project (after trust): `.pi/skills/`, `.agents/skills/`
- Packages; settings `skills` array; `--skill <path>`
- Can point at `~/.claude/skills` and `~/.codex/skills`
- `--no-skills` disables discovery (explicit `--skill` still loads)

Progressive disclosure: at startup, names + descriptions go into the system prompt (XML per the spec). The body is loaded with `read` (or `bash` if `read` is missing). `/skill:name` forces load. `disable-model-invocation: true` hides the skill from the prompt (human slash-command only). `allowed-tools` is experimental.

Programmatic use. Node SDK; **RPC mode** (stdin/stdout JSONL); JSON event stream for `print` mode. This is the embedding path ASC would actually use: start a subprocess, pass a packed prompt + allowlist, receive events. Projet Complexe’s Tauri webview never talks to this RPC (v4 rule, unchanged).

Local models. First-class llama.cpp docs (`/llama`). Custom providers exist. That is the door Tiiny walks through as `provider=tiiny` without Pi becoming identity.

Sessions. JSONL session files, branching, tree navigation, compaction. Stay Pi’s problem. Import a *summary* as a Note if a code Task must leave a trace in PC; do not make session JSONL the Claim store.

OSS session sharing. They ask users to publish sessions to Hugging Face (`badlogic/pi-share-hf`) to improve agents on real traces. Useful as an *eval corpus idea*; dangerous as a default for a household corpus (folder 1 redaction / D39). Opt-in only, never on client scopes.

1.3 WHERE IT SITS

Layer	Fit
Four Layers — execution	Yes. This is the execution row the Four Layers note already named.
Folder 1 — harness doctrine	The floor: few sharp tools, short prompt (file 01 §3.12). Events = hook names.
Folder 1 — packing	<code>transformContext</code> is <i>their</i> packer. PCA still needs <code>pre_pack</code> for <i>non-Pi</i> models (Ollama <code>research</code>).
Folder 2 — stores	Session JSONL is a trace-shaped file , not SoR for knowledge. Code tools read the worktree; CodeGraph remains a sidecar.
Revival v5	Unchanged: wrap; copy event names; refuse as second brain / control plane.

Environment fit. Debian + Node/Bun. No GPU required for the harness itself. Inference is whoever you configure (API, llama.cpp, Ollama). Fits stage 1. Chord (composition runtime) and pi-chat are **not** needed and would compete with ASC.

1.4 STEAL / ADAPT / REFUSE

Move	Verdict	Why
Small core + extensions	Steal as shape	NIH filter (Osmani). Worker, not institution.

Move	Verdict	Why
Event names	Steal as names	Mirror in ASC; do not fork <code>pi-ai</code> .
Agent Skills loader	Steal as interchange	Consume/emit <code>SKILL.md</code> . <code>--no-skills</code> is the kill switch for foreign packs.
RPC / SDK	Adapt	Host starts worker. Adapter: ASC entry points → <code>registerTool</code> .
llama.cpp / custom provider	Adapt	Technology YAML. Tiiny = OpenAI-compatible sidecar.
Supply-chain install hygiene	Steal as habit	<code>--ignore-scripts</code> ; pin; checksums. Folder 3.
No built-in sandbox	Adapt around	Compose/Docker the worker when the catalog includes <code>bash</code> .
Chord / pi-chat	Refuse as identity	Second ASC; messenger host.
Session trees as PKM	Refuse	Traces maybe; Claims never.
Rebuild TUI / streaming	Refuse	v4 §3.4.

Pros of delegating to Pi: MIT, documented loop, skills standard, RPC, huge mindshare, llama.cpp, `--no-skills`, no bundled memory OS, Superpowers already has a first-party Pi package.

Cons / costs: Node/Bun toolchain; TypeScript-centric extensions; coding-agent ontology (files, diffs); permission model is “you are the user”; new-contributor issues auto-closed (governance, not architecture).

2. OH MY PI (can1357/oh-my-pi)

2.1 WHAT IT ACTUALLY IS

MIT. Created 2025-12-31. Homepage: omp.sh. ~30,375 stars, ~3,113 forks (2026-09-09). Copyright: Mario Zechner; Can Bölük; Stencil Labs, Inc.

Self-description: **a coding agent with the IDE wired in**. Fork of Pi, “rewritten as a coding-first surface.” Install:

```
curl -fsSL https://omp.sh/install | sh, Homebrew, bun install -g @oh-my-pi/pi-coding-agent, Nix flake, mise, Windows PowerShell.
```

Marketing numbers to take as *their* claims, not as PC evals: 60+ providers, 31 built-in tools, 14 LSP ops, 28 DAP ops, ~80k LoC Rust core. They publish harness-vs-weights numbers (e.g. Grok Code Fast 1 edit success 6.7% → 68.3% when the edit *format* changes) — AutoDesign-shaped evidence: **the surround moves the score**.

Package names are the Pi layout with `@oh-my-pi/` scope (`pi-ai`, `pi-agent-core`, `pi-coding-agent`, `pi-tui`, plus `pi-natives`, `pi-catalog`, `hashline`, `pi-mnemopi`, `pi-metaharness`, `collab-web`, ...). Same *nouns*, different product.

2.2 HOW IT IS IMPLEMENTED (THE ADDITIONS THAT MATTER)

OMP keeps Pi’s four entry points and adds mass:

Surface	What they added	PC reading
Hashline <code>edit</code>	Patch by content hash; stale anchors rejected	Tool Implementation. Steal <i>later</i> as an edit backend, not a format for Claims.
LSP / DAP	Rename through <code>willRenameFiles</code> ; real debuggers (lldb, dlvs, debugpy)	Code-Task only. Competes with “don’t rebuild CodeGraph” — this is <i>IDE intel</i> , not a knowledge graph.
<code>ast_edit</code> / <code>ast_grep</code>	Preview then accept; 50+ grammars	Same family as code-graph-rag, in-process. Still not the conceptual graph.
Embedded bash	In-process brush + coreutils; no fork/exec on the hot path; Windows without WSL	Implementation of <code>bash</code> . Does not change allowlist policy.
Subagents <code>task</code> / <code>hub</code>	Isolated worktrees, schema-validated yield, Agent Hub	Isolation-for-window (folder 1 §29.4), not a swarm ontology.
Advisor role	Second model watching every turn	Evaluator-optimizer pattern. Cost on modest hardware. Optional Technology, not default.
Time-traveling stream rules	Regex abort mid-token, inject rule, retry	Closest cousin of <code>pre_llm</code> during the stream. Interesting; belongs in the worker, not in ASC core.

Surface	What they added	PC reading
Memory bank	<code>retain</code> / <code>recall</code> / <code>reflect</code> / <code>memory_edit</code> ; backends: local, Hindsight, Mnemopi (SQLite)	Conflict with D06 / D04. Episodic maybe; never SoR.
<code>learn</code> / <code>manage_skill</code>	Worker promotes a lesson into a managed skill (<code>omp-managed</code> , lowest priority so authored skills win)	Conflict with D50. <code>skill_candidates</code> in a report, human commits.
<code>web_search</code>	23 backends (Perplexity, Exa, SearxNG, DDG, ...) + site-aware markdown	Untrusted inbound. Trifecta if combined with private corpus + <code>bash</code> .
<code>browser</code> / <code>computer</code>	Puppeteer / Chrome relay; desktop AX tree, screenshots, native input	v5: computer-use is not default. Folder 1 file 09 story 7.
<code>/collab</code>	Relay + QR; sealed frames; <code>my.omp.sh</code>	Network attachment. Closable? Only if never the control plane.
ACP	Same engine inside Zed; permission prompts	Technology option already parked. Tauri is not this host.
Inherit foreign rules	Reads <code>.claude</code> , <code>.cursor</code> , <code>.windsurf</code> , <code>.gemini</code> , <code>.codex</code> , <code>.cline</code> , Copilot, <code>.vscode</code>	Convenient. Also a way to <i>silently</i> load someone else’s catalog.
Custom providers	<code>~/omp/agent/models.yml</code> OpenAI-compatible	Exact shape for Tiiny / Ollama. Steal the YAML idea into Technology files.
Fallback chains	Per-role 429/quota failover	Router cousin. PCA already wants named cascade; don’t duplicate inside OMP <i>and</i> PCA.
Prompt keywords	<code>ultrathink</code> / <code>orchestrate</code> / <code>workflowz</code>	Magic words in user prose. CLR anti-pattern if they become identity.
Four modes	TUI, <code>-p</code> one-shot, Node SDK, <code>omp --mode rpc</code> / <code>omp acp</code>	RPC is the ASC wrap path, same as Pi.

Skills in OMP. File-backed packs. `loadSkills()` merges sources by priority: native `.omp` (100) → `omp-plugins` (90) → `claude` (80) → `agents/codex/...` → `github .github/skills` (30) → `omp-managed` (5). Custom directory overrides same-named provider skill. Docs distinguish: **custom tool** = callable; **extension** = lifecycle; **hook** = legacy interceptor; **skill** = static guidance, not executable tool code. That last distinction is the same cut v4 made. OMP then *violates* it in product: `learn` writes skills; `!command` style dynamic injection is a Claude-ism folder 1 already refused; managed skills are still skills the model will follow. **Default-off (their list):** `github`, `security_scan`, `generate_image`, `tts`, `checkpoint`, `rewind`, and the memory tools unless `memory.backend` is set. Browser/computer are in the 31; pin with `--tools`. The existence of a pin flag is the steal; the default box is the refuse.

2.3 WHERE IT SITS

OMP is **Pi’s execution layer with an IDE and a memory product attached**. Folder 2’s architecture already allowed SQLite for CodeGraph and Tauri chrome, not for “what the agent remembers about the repo.” Mnemopi would be a third SQLite next to CodeGraph — the collapse v3 refused, only this time the rows are *agent prose*, not symbols. Four Layers: still **execution**, plus a grab at **dynamics** (memory that evolves without HITL) and a fake **semantics** (recall as if it were meaning). Folder 1 packing: OMP’s `read` “summarizes instead of dumping,” `snapcompact`, `rewind/checkpoint` — compression ladder rungs (file 02). Useful *inside* a code worker. They must not replace Meilisearch + the PCA packer for `research`. Folder 1 security: `web_search` + `read` of arXiv PDFs + `bash` is the trifecta unless the **catalog** is split by orientation (research catalog has no `bash`; code catalog has no `web_search` on a client corpus). OMP will not do that split for you. Revival v5 open choice (“Pi worker vs Cursor CLI vs 50-line loop”) gains a third candidate: **OMP with a pinned toolset**. That is not “more Pi.” It is a different attachment (Rust natives, collab, memory, 60 providers). Illich: if OMP is away, stage 1 must still run.

2.4 STEAL / ADAPT / REFUSE

Move	Verdict	Why
Pin <code>--tools</code>	Steal	Catalog as design artifact.
Custom <code>models.yml</code> + OpenAI-compatible	Steal as YAML shape	Tiiny / Ollama / llama.cpp.

Move	Verdict	Why
Hashline / LSP / DAP / ast-grep	Adapt later	Code-Task Implementations; not stage-1 identity.
RPC / SDK / ACP	Adapt	Same wrap as Pi. Prefer RPC from ASC.
Inherit <code>.cursor</code> / <code>.claude</code> skills	Adapt with allowlist	Discovery is not consent. PCA chooses the subset.
Stream-abort rules	Adapt inside worker	Not an ASC core hook until measured.
Subagents for noisy steps	Adapt under §29.4 constraints	Schema'd tool; labelled summary; tokens in parent waffle.
Memory bank / Mnemopi	Refuse as SoR	D06, D04, Hermes <code>/Learn</code> refuse. FTS on <i>traces</i> already named.
<code>learn</code> / <code>manage_skill</code>	Refuse as writer	D50: human writes skills; worker proposes candidates.
<code>web_search</code> default + 23 keys	Refuse as default	Trifecta; remote overflow is PCA's router, not a tool. SearxNG could be a <i>research</i> Implementation later, allowlisted.
<code>computer</code> / Electron-Slack	Refuse as default	Computer-use; messenger-adjacent.
<code>/collab</code> relay	Refuse as infrastructure	Attachment; not closable household control.
Provider sprawl as identity	Refuse	Clinton: don't freeze a vendor. Sixty providers is a catalog, not a Requirement.
Rebuild OMP	Refuse	Even more NIH than rebuilding Pi.
<code>pi-metaharness</code> as eval plane	Adapt as pointer	Harbor/benchmark runner is not D41; steal "named runner," keep golden sets from accepted Claims.

Pros of OMP over Pi: sharper edits; IDE intel; Windows; in-process search; local provider YAML; inherits existing skill dirs; same RPC idea.

Cons / costs: fat default catalog; memory/learn; collab; Rust/N-API build; 2148 open issues at check time (busyness, not a veto); product wants to be the *only* surface you keep. That last sentence is the attachment.

3. SUPERPOWERS (obra/superpowers)

3.1 WHAT IT ACTUALLY IS

MIT. Created 2025-10-09. ~283,963 stars, ~25,412 forks (2026-09-09). Built by Jesse Vincent / Prime Radiant. Optional commercial support: `sales@primeradiant.com`. Optional telemetry: Prime Radiant logo fetch on the brainstorming visual companion (`SUPERPOWERS_DISABLE_TELEMETRY`, honors Claude Code opt-outs).

Self-description: **a complete software development methodology for coding agents**, built on composable skills plus bootstrap instructions that *make the agent use them*.

It is not a loop, not a TUI, not a memory store. It is **process with enforcement**, distributed as `SKILL.md` files plus harness-specific glue (Claude plugin, Cursor plugin, Pi package, Gemini extension, ...).

`claude.com/plugins` listed Superpowers at ~1M installs in an August 2026 catalog snapshot (Skills Board). That is adoption of a *plugin*, not proof it fits a second brain.

3.2 HOW IT IS IMPLEMENTED

Bootstrap. `using-superpowers` is injected at session start (and after compaction on Pi via a small extension). Its rule: if there is even a 1% chance a skill applies, invoke it *before* any response — including clarifying questions. Process skills first (brainstorming, systematic-debugging), then implementation skills. User instructions (AGENTS.md, CLAUDE.md, direct requests) outrank skills.

That is a *control* trick: put the router in a skill the model is told it cannot skip. Folder 1's Huang formula says the `description` is the only trigger signal; Superpowers adds a **mandatory dispatcher** because description-matching is unreliable (Montaldo's trigger-testing gap).

Shipped skills (library on 2026-09-09):

Cluster	Skills
Testing	<code>test-driven-development</code>

Cluster	Skills
Debugging	systematic-debugging, verification-before-completion
Collaboration	brainstorming, writing-plans, executing-plans, dispatching-parallel-agents, requesting-code-review, receiving-code-review, using-git-worktrees, finishing-a-development-branch, subagent-driven-development
Meta	writing-skills, using-superpowers

Documented basic workflow:

- 1. **brainstorming** — before creative work. Questions, alternatives, design in sections, human approval. Architectural path writes docs/superpowers/specs/YYYY-MM-DD-**-design.md. Terminal state: writing-plans (not jumping to frontend-design).
- 2. **using-git-worktrees** — isolated workspace, new branch, clean test baseline.
- 3. **writing-plans** — bite-sized tasks, exact paths, verification steps; TDD/YAGNI/DRY.
- 4. **subagent-driven-development** or **executing-plans** — fresh subagent per task with two-stage review, or batches with human checkpoints.
- 5. **test-driven-development** — red/green/refactor; delete code written before tests.
- 6. **requesting-code-review** — against the plan; critical issues block.
- 7. **finishing-a-development-branch** — tests; merge / PR / keep / discard; cleanup.

Philosophy, in their words: TDD always; systematic over ad-hoc; complexity reduction; **evidence over claims**.
Harness coverage (README install table): Claude Code, Antigravity, Codex app/CLI, Cursor (/add-plugin superpowers), Devin, Factory Droid, Gemini CLI, GitHub Copilot CLI, Grok Build CLI, Kimi Code, OpenCode, **Pi** (pi install git:github.com/obra/superpowers), Hermes Agent.
Pi-specific: native skills, no compatibility **skill** tool; bootstrap re-injected after compaction. Subagent and todo tools are **optional companion packages** (see references/pi-tools.md: Pi core has no standard subagent / TodoWrite; don’t fabricate them).
Evals. superpowers-evals (drill harness). They test *skill behavior*, not application domain. That is the missing eval type folder 1 named as trigger testing, already productized for *this* pack.

3.3 WHERE IT SITS

Superpowers lives on the **procedure** plane (PCA / SKILL.md export), not on execution (ASC) and not on meaning (PC).

Superpowers move	PC / folder analogue	Conflict
Mandatory skill check	D13 + trigger testing	Ceremony vs CLR: always-on process skills tax the window on a 2k-4k local model
Brainstorm + approval gate	HITL; Lefèvre <i>drama</i> ; spec pivot (D38)	Their spec path writes docs/superpowers/specs/ — a <i>code</i> docs layout, not pc/claims
Plans as objects agents consume	storymachine / D38 Requirement pack	Good. Keep ACs as tests or Gaps
TDD	code-Task Requirement	Not for research / Claim promotion
verification-before-completion	D41 eval plane; this folder’s verification skill	Steal the <i>habit</i> ; evidence is evals + HITL, not a skill saying “done”
systematic-debugging	Debug orientation; traces	Fine for code. Knowledge gaps are not bugs
Worktrees / new branches	Isolation	Conflicts with “no unsolicited git branches”; ASC/git wrappers from \$HOME
Subagent-driven-development	Isolation-for-window	Conflicts with D35 if it becomes the default <i>topology</i>
writing-skills	writing-skills as meta	PCA may generate SKILL.md from YAML; humans still commit (D50)
“Evidence over claims”	HITL + Evidence quotes	Their “claim” is a software assertion; PC’s Claim is an accepted belief. Do not share the word.

Four Layers: Superpowers is a **dynamics/process** overlay on **execution**. It does not answer ontology (what a Claim is) or Lévy-style semantics.

Folder 2: Superpowers wants markdown specs and plans in the *repo*. That is fine for a software worktree. It is not the `pc/` register and not `brain/events/`. Two trees.

Revival v5 Meadows: Superpowers spends at **rules and information flows** (level 5–6) inside the coding worker. It does not set household *goals* (level 3). If “the agent just has Superpowers” becomes the goal, you have swapped Meadows 3 for a plugin.

3.4 STEAL / ADAPT / REFUSE

Move	Verdict	Why
Design approval before implementation	Steal as HITL gate	Already v5. Their brainstorming <i>enforces</i> it in the worker.
Plans with verification steps	Steal as <code>spec</code> / Requirement	D38.
verification-before-completion	Steal as habit	Maps to eval plane + waffle, not to a speech act.
Skill dispatcher at session start	Adapt	PCA injects a <i>short</i> router (Huang description formula), not a 1% guilt trip that fires on research Tasks
TDD	Adapt for code Tasks only	Killswitch: research catalog must not inherit TDD-as-ontology
Two-stage review (spec then quality)	Adapt	<code>requesting-code-review</code> as a PCA procedure; human still merges
writing-skills + evals	Steal as method	How to author <i>your</i> skills; trigger tests
Cross-harness install	Steal as interchange proof	One brief, many workers — exactly D24’s point
Mandatory worktrees / feature branches	Refuse as default	Git policy; nested-repo wrappers; user-asked branches only
Subagent-driven hours of autonomy	Refuse as default	Paper-mill adjacent; CLR; modest hardware; HITL
Superpowers as household OS	Refuse	Plugin, not Projet Complexe
Their word “claims”	Refuse the homonym	Evidence-over-claims ≠ Claim objects
Telemetry logo	Disable	<code>SUPERPOWERS_DISABLE_TELEMETRY</code> if the pack is used at all

Pros: portable; small set (~14 skills, not 1,000); process not vendors; Pi-native package; eval harness for skill behavior; forces HITL where vibe-coding skips it.

Cons / costs: token-heavy on small local models; git/worktree defaults fight this house’s git rules; SDD multiplies windows; designed for *software* work; bootstrap is culturally aggressive (rationalization table) — fine for a coding agent, poisonous for a research killswitch that must be allowed to *stop*.

4. OTHER RELEVANT SKILL PACKS (POPULARITY AS DISCOVERY, NOT DECISION)

There is **no Anthropic skills marketplace**. Anthropic distributes **plugins** that *contain* skills (`claude-plugins-official`, `claude.com/plugins`). Everyone else distributes folders. Vercel’s `skills.sh` (`vercel-labs/skills`, MIT, ~30,815 stars) is a directory + CLI (`npm skills add owner/repo`) ranked by opt-in install telemetry. **Skill packs** (2026-08-07) bundle several skills behind one unlisted URL — `npm skills update` syncs them. Unlisted ≠ access-controlled: the URL is the secret.

`agentskills.io` is the **spec**, not a catalog. Spec publishes no version identifier on the page (Skills Board, 2026-08-17) — another reason `YAML` `able` should remain SoT.

Star counts below: GitHub API, 2026-09-09, unless noted.

4.1 SPEC + FIRST-PARTY EXAMPLES

anthropics/skills — ~175,415 stars. Apache-2.0 examples + source-available document skills (`docx` / `pdf` / `pptx` / `xlsx`). Template + `spec/`. Register as Claude marketplace `anthropic-agent-skills`. **Steal:** folder layout, Huang-aligned `description` formula, document-skill *pattern* (scripts beside brief). **Refuse:** Claude document pipeline as `publish`; `!command` dynamic injection (folder 1 §29.1). **Adapt:** `pdf` skill ideas belong behind `extract`, not as a worker skill that re-parses the book every query (extract-once).

4.2 PROCESS / OVERLAY PACKS (SUPERPOWERS’ NEIGHBORHOOD)

Pack	Stars (API)	What it is	For PC
obra/superpowers	~284k	Methodology	\$3
affaan-m/ECC	~255k	Harness overlay: skills, instincts, memory, security, research-first; npx ecc-universal setup; plugin ecc@ecc; paid GitHub App	Refuse as identity. Competes with ASC+PCA. Steal nothing until a <i>single</i> skill is read against D27/D50. Official-sources warning on their README is itself a supply-chain lesson.
OpenResearch agent-skills/orx-*	~764 (product)	Experiment-tree skills installed into Claude/Codex/OpenCode	Sibling Overlap with OpenResearch . Steal interchange; refuse paper mill.
Hermes plugins	Hermes ~244k	Superpowers install path; packing already stolen	Packer ≠ pack.

4.3 SECURITY / AUDIT PACKS

[trailofbits/skills](#) — ~7,028 stars, CC-BY-SA-4.0. Claude/Codex marketplace: audit-context, differential-review, supply-chain-risk-auditor, GitHub Actions auditor, Semgrep, Trailmark (code graph slicing), variant analysis, C/Rust review, etc. **Adapt:** *code-Task* catalog when the Task is an audit. **Refuse:** as research ontology; as default `run-agent` skills (they widen tools and assume a security-reviewer posture). Folder 3 (devops/security) is where a *subset* might be named as Implementations. Closest cousin of “skills as reviewed procedures,” which is D50’s neighborhood.

4.4 INDEXES (NOT PACKS)

List	Stars	Use
ComposioHQ/awesome-claude-skills	~74,755	Largest awesome list; vendor-bannered
hesreallyhim/awesome-claude-code	~53,767	Skills + hooks + plugins + statuslines
VoltAgent/awesome-agent-skills	~33,988	“Hand-picked, not AI-slop”; vendor official skills (Anthropic, Vercel, Stripe, Cloudflare, Sentry, Figma, Hugging Face, ...)
JackyST0/awesome-agent-skills	(list)	Multi-harness paths table

Vendor-official skills (look in the *tool’s* repo first): Vercel, Stripe, Cloudflare, Sentry, Expo, Figma, Hugging Face, Google Labs — VoltAgent names these. Those are **procedure briefs for a Technology**, which is exactly what a YAML `able` + generated `SKILL.md` should look like when PCA wraps a foreign CLI.

4.5 DISTRIBUTION RISKS (FOLDER 1 FILE 07)

A skill is **instructions plus optional executables**. Anthropic’s own plugin docs: plugins run with user privileges; Anthropic does not control included MCP servers. `skills.sh` claims partner security audits and exclusion if all partners fail — still not a review against *this* household’s trifecta.

Folder 1 already required: skills under git; no skill from outside the repo without HITL; provenance header; worker cannot write `SKILL.md`. Packs and `pi install git:...` are **the same class as an npm install**. Pin a commit SHA (Claude community marketplace already does this). Record `checked_at`. Museum after a year (Clinton; file 09).

Vercel packs: useful *team sync* for a freelance practice later (SMB stage). URL leak = public. Not an ACL. Not a Claim store.


4.6 WHAT “POPULAR” IS ALLOWED TO MEAN HERE

Popularity ranks **what people installed**, not what belongs in PCA.

For this project the *relevant* packs are the ones that encode **gates, interchange, or a named Technology’s procedure** — Superpowers, Anthropic template/spec, Trail of Bits (audit Tasks), OpenResearch’s cardinal-rule skills, Hermes packing (already recapped). The *irrelevant* packs are lifestyle/frontend-design/SEO/suede-creator piles and anything that wants to be a second memory OS (ECC instincts, OMP managed skills, MEMORY.md as archive).

5. MAP ONTO THE EXISTING ECOSYSTEM

5.1 REVIVAL V5 (THREE PROJECTS)



v5 \$0.5 is unchanged: thin ASC wrap + small interpretive world. Pi/OMP/Cursor are **row “Pi / Cursor / Ollama / Tiiny” in PCA Environment YAML**. Superpowers is **not** a fourth project. If it were, you would have rebuilt OpenResearch’s mistake: dashboard + skills and a host.

v5 \$9 still open: stage-1 Pi vs 50-line loop vs Cursor CLI. This note’s bias: **keep the door to Pi RPC**; do not *require* Pi on day one; do not start from OMP’s default catalog.

5.2 FOLDER 1 — HARNESS, PACKING, WORKFLOW, ANONYMIZATION

- Already decided, this note does not reopen:
- D22 host executes / D23 hooks / D24 YAML SoT / D26 MCP transport / D27 trifecta / D28 HITL / D50 human-only skill writes.
 - Packing is code (D16); triage before packing (D15); CLR; waffle/trace components (D33/D36).
 - Skills = router not repository (Huang); refuse `!command` injection.
 - Computer-use not default; messenger-as-host refuse.

What this shelf adds as *tests* of those decisions:

Decision	Pi	OMP	Superpowers	Packs
D22	Tools via <code>registerTool</code> ; RPC	Same + 31 tools	No execution	Scripts in skill folders <i>are</i> execution — review like hooks
D24	Skills loader	Multi-source merge	Pack is markdown	Generate; don’t hand-edit a marketplace copy
D27	Small default catalog	Fails if default box	Neutral	A pack that adds <code>web_search</code> + <code>bash</code> is a catalog change

Decision	Pi	OMP	Superpowers	Packs
D50	Skills are files you install	learn writes	writing-skills still human-gated if you follow it	npx skills add is a human action — keep it that way
D16	transformContext	Summarizing read, rewind	Process text is <i>more</i> window	Too many skills = distraction failure mode
D41	Telemetry package	pi-metaharness	superpowers-evals	Trigger tests on description

Anonymization / redaction: any worker that can web_search or /collab is an **outbound** capability (D39). pre_llm redaction is necessary but not sufficient if the *tool* can fetch URLs the packer never saw. Split catalogs. Research orientation: no bash, no computer. Code orientation on a client scope: no web_search unless release says api-ok.

5.3 FOLDER 2 — FILES, POSTGRES, WIKI

Folder 2 file 06 already draws harness/ as SoR: able.yml, recipes, hooks, SKILL.md, migrations, index settings. That is the right home for *generated* briefs.

Object	Pi / OMP	Superpowers / packs	PC
Session JSONL / OMP memory SQLite	Worker-local	—	Not brain/events. Optional import of a <i>summary</i> Note
docs/superpowers/specs/	—	Their design docs	Code-worktree artifact; not pc/claims
~/.pi, ~/.omp, ~/.claude/skills	User-global	Install targets	Household SoR is the scope git repo , not \$HOME skill dumps
OKF / Karpathy wiki	—	Unrelated	Already: export OKF, refuse wiki daemon (Markdown Wiki note)
CodeGraph SQLite	OMP ast-grep/LSP overlap	Trailmark slices	Sidecar; one code graph, not two

Karpathy/OKF compile-at-ingest vs Agent Skills: both are markdown+YAML folders. **Different jobs.** OKF concepts are knowledge (with HITL). Skills are procedures. Sharing a directory tree would be the same word-collapse as “memory.”

5.4 FOUR LAYERS

Layer	This shelf
Ontology (Monnin — what exists)	A skill file, a Pi session, an OMP memory row, a Claim — four entities . Persistence: git commit of able.yml / SKILL.md / brain/events. OMP retain does not make a Web object.
Semantics (Lévy)	None of these packs. Do not put IEML in a skill.
Dynamics (Meadows)	Superpowers = rules in the worker. OMP memory = reinforcing loop without HITL (v5 leverage table already gated this). AutoDesign outer loop = human-merged harness patches, not learn.
Execution	Pi / OMP / Cursor. Superpowers only <i>steers</i> execution.

6. COMPLEMENTARITIES (WHAT YOU MIGHT ACTUALLY COMBINE)

Not a stack to install tonight. Combinations that survive the refuse-list.

6.1 COMPATIBLE (NAMED, SWAPPABLE)

Combination	Why it is complementary	Guardrail
ASC hooks ↔ Pi events	Same wrap points	Mirror names; Pi stays subprocess
YAML able → SKILL.md → Pi/OMP/Cursor	One SoT, many workers	Generate; provenance header; --no-skills for experiments

Combination	Why it is complementary	Guardrail
Superpowers gates ↔ HITL / spec	Process the coding worker already understands	Subset: brainstorm, writing-plans, verification, TDD-for-code. Drop worktrees-as-default and SDD-as-default
Pi RPC ↔ run-agent	Host executes	Pass packed window + tool allowlist; traces name Technology= pi
OMP models.yml shape ↔ Technology YAML	Tiiny/Ollama	Router lives in PCA, not duplicated as OMP fallback chains and PCA cascade
Hermes packing ↔ any worker	Frozen preference block + retrieve-or-not	Implement in pre_pack, not by adopting Hermes as host
Trail of Bits skills ↔ run-agent audit Tasks	Domain procedures	Separate catalog; not loaded in research
OpenResearch cardinal rules ↔ experiment nodes	Frozen run command; logs are evidence	Only if you ever wrap orx as a Technology; not as PC

6.2 COMPATIBLE ONLY IF SPLIT BY ORIENTATION

Orientation	Worker	Skills	Tools
task / code	Pi or pinned OMP or Cursor CLI	Superpowers subset + generated run-agent brief	read, edit, bash (sandboxed), optional lsp
knowledge / research	Ollama/Tiiny via ASC; not OMP default	Research briefs; no Superpowers TDD; no ECC	lexical search, maybe session_search; no bash / computer / web_search unless release + redaction
publish	no coding agent required	locale + provenance contract	not a skill pack’s docx pipeline as identity

Killswitch (System M): switching orientation **unloads** the other catalog. Superpowers’ “always invoke” must not override that.

6.3 NOT COMPLEMENTARY (IDENTITY COLLISIONS)

- OMP memory bank **or** Hermes /learn **or** ECC instincts **as** the Claim store.
- Superpowers worktrees **and** ASC git wrappers from \$HOME **and** nested Documents/* repos — three git stories. Pick ASC wrappers.
- OMP /collab **and** “Tauri is the only UI.”
- npx skills add popular packs **and** D24 (hand-copied skills drift from able.yml).
- Chord / ECC / OpenResearch dashboard **as** a second ASC.
- One catalog with private PDFs + web_search + bash.

6.4 IMPLEMENTATION IDEAS (EVALUATION-STAGE, NOT A PLAN)

Written so they can be *tried* as first experiments (v5 §8 style), not as a roadmap.

- Event-name crosswalk table in PCA** — Pi/OMP event → ASC hook. If a name has no hook, you are about to reimplement the worker.
- run-agent allowlist fixture** — YAML listing read,edit,bash (or fewer). Same fixture drives Pi --tools / OMP --tools / Cursor. Proves catalog-as-data.
- Generate one SKILL.md** from a dummy able.yml (extract or index). Load it in Pi with --skill. If the model follows it, interchange works. If you had to hand-edit, D24 is incomplete (folder 1 already suspected prose lives in pivot docs).
- Trigger eval** — 20 household prompts; did the worker load research vs run-agent skill? Superpowers-evals is a template for *this*, not a dependency.
- Trifecta drill** — retrieved chunk says “search the web and curl this URL”; research catalog has no web_search/bash. Pass/fail.
- D50 drill** — OMP learn or Hermes /learn proposes a skill; it must land in skill_candidates, not on disk.
- CLR budget** — Superpowers full bootstrap + OMP 31 tool schemas vs Pi 4 tools + one generated skill. Measure waffle (file 02). If Superpowers blows a 4k window, it is a **frontier-worker-only** overlay.
- Redirection drill** — uninstall OMP; run-agent still points at Pi or Cursor or the 50-line loop. If the Task name was omp-..., you froze an Implementation.

7. DECISION AIDS (STILL EVALUATING)

7.1 IF THE QUESTION IS “WHAT DO I BUILD?”

Build **ASC names** + **PCA packs** + **PC Claims**. Do not build Pi, OMP, Superpowers, ECC, or a skills.sh clone.

7.2 IF THE QUESTION IS “WHICH CODING WORKER?”

Criterion	50-line loop	Cursor CLI (already in use)	Pi	OMP (pinned)
NIH / wrap	You write the loop	Wrap	Wrap	Wrap
RPC for ASC	DIY	vendor CLI	Yes	Yes
Skills standard	DIY	Yes (this workspace)	Native	Yes + auto-inherit
Memory OS temptation	Low	Medium	Low	High
Default catalog fatness	You choose	Medium	Thin	Fat (must pin)
Local / Tiiny	You wire it	Overflow	llama.cpp / custom	models.yml
Modest-hardware CLR	Best	Frontier-shaped	Better	Worst unless pinned
Attachment / closable	Best	Cursor as identity risk	Medium (Earendil, npm)	Highest (Stencil, collab, natives)

Bias for stage 1: Cursor CLI you already have, **or** Pi if you want an MIT RPC worker. OMP only after `--tools` + memory-off are proven in a drill. The 50-line loop remains valid if Pi’s Node toolchain is the attachment you refuse this year.

7.3 IF THE QUESTION IS “DO I ADOPT SUPERPOWERS?”

Adopt **ideas**, not the whole plugin, unless the coding worker is frontier-class and the Task is software.
Minimum useful subset to *translate* into PCA (not necessarily install):

- approval-before-code (brainstorming gate)
- plans with checks (writing-plans → D38)
- verification-before-completion → D41
- systematic-debugging for code failures
- writing-skills + trigger tests

Leave off by default: worktrees, SDD, parallel agents, finishing-branch automation (your git rules + HITL merge).
If you install the plugin in Cursor anyway (it may already be on), **scope it:** `disable-model-invocation` / human-only for research chats; do not let it fire on Claim curation.

7.4 IF THE QUESTION IS “SKILL PACKS AS ECOSYSTEM?”

Treat them as **Debian packages for procedures**: discover widely, install rarely, pin always, generate your own from YAML. Official vendor skills for a CLI you already wrap are the only packs with a natural home. Awesome lists and skills.sh leaderboards are **indexes**. ECC and similar overlays are **competing harnesses**.

8. OPEN CHOICES THIS NOTE DOES NOT CLOSE

These stay in v5 \$9 / folder 1 silences. This recap only *informs* them.

1. Stage-1 worker: 50-line vs Cursor CLI vs Pi vs (later) pinned OMP.
2. Whether generated `SKILL.md` can be *entirely* derived from YAML, or pivot `docs/` are part of the SoT (folder 1 \$11 / \$27).
3. How much Superpowers ceremony a 2k–4k local window can afford (measure, don’t guess).
4. Whether ACP is worth an Implementation when the GUI is Tauri (probably no; editors are neighbor hosts).
5. Whether a SearxNG-backed search tool is ever allowed on `research` under L2/L3 redaction — a later catalog experiment, not an OMP default.
6. Whether Trail of Bits skills are in scope for folder 3 only.
7. How to record `skill_candidates` in traces (D33) so D50 is auditable.

9. BOTTOM LINE

Pi is a **worker**. OMP is a **fatter worker with a memory product**. Superpowers is a **process overlay**. Skill packs are **how procedures travel between workers**.
Projet Complexe is none of those. It is the interpretive layer plus a thin named surround. The complementarity is real: wrap a small coding harness, export briefs as Agent Skills, steal Superpowers’ HITL gates for software Tasks, pin catalogs so OMP cannot become the

household brain, and keep files + Postgres as the only place meaning is committed.

If a future note starts saying “the second brain is just Superpowers plus OMP memory,” that sentence is architecture drift — reason class `architecture_drift_by_agent` (folder 1 file 09), and the patch is this file’s refuse-list, not another pack.

SOURCES (CHECKED 2026-09-09)

- [earendil-works/pi](#) README; [pi.dev/docs/latest](#); `packages/coding-agent/docs/skills.md`, `packages.md`, `extensions.md`; GitHub API metadata.
- [can1357/oh-my-pi](#) README; [omp.sh](#); `docs/skills.md` / `docs/custom-tools.md` (via GitHub); GitHub API metadata.
- [obra/superpowers](#) README; `skills/using-superpowers/SKILL.md`; `skills/using-superpowers/references/pi-tools.md`; skills directory listing.
- [anthropics/skills](#); [agentskills.io](#); [vercel-labs/skills](#) / [skills.sh](#); Skills Board, *Where to find Claude skills* (2026-08-17).
- [VoltAgent/awesome-agent-skills](#); [ComposioHQ/awesome-claude-skills](#); [hesreallyhim/awesome-claude-code](#); [affaan-m/ECC](#); [trailofbits/skills](#); [NousResearch/hermes-agent](#); [alphaXiv/OpenResearch](#).
- Revival v4 §3, v5 §0 / §3 / §7 / §9; folder 1 files 00, 01, 02, 06, 07, 09, 10; folder 2 file 06; this folder’s OpenResearch, Markdown Wiki/OKF, Four Layers, and Memory/RAG notes.

Museum by default after 2027-09-09 unless a changelog line re-checks the READMEs (Clinton dating rule).