

PROJET COMPLEXE 2026 REVIVAL (V4)

HOOKS AROUND LLM CALLS, TOOLS AS ASC ENTRY POINTS, DSL VERSUS MCP

Date: 2026-08-21

Status: architecture note / design instrument (not a spec, not an implementation plan)

Keeps: Revival v2 (three-project cut, killswitch, genericity scale) and Revival v3 (Postgres + pgvector + Meiliseach as retrieval; extract as a capability; RAG as context regulation).

Adds: (1) a decision on *how agents get tools* — MCP vs JSON/Python registries vs Claude/Pi “skills” vs ASC hooks + entry points + DSL; (2) after the full shelf, a second cut: language, routing, curated knowledge, and which doors not to close. The first cut is the original question. The second is what that question is *for*.

This note exists because a theoretically attractive idea is easy to over-build:


Wrap every LLM call in ASC hooks that pre-process and post-process the prompt. Implement what the industry calls tools, actions, or skills as ordinary ASC entry points. Drop MCP. Use DSL instead of JSON and Python.

The idea is **not wrong**. Parts of it are already how ASC works. Parts of it are already how every serious agent *harness* works in 2025–2026. The useful question is not “is this possible?” It is:

Which of those layers is **ASC-generic** (worth putting in core, useful beyond one personal knowledge-management app), which belongs in **Projet Complexe ASC** as a named Implementation, and which should be **delegated** to a project that already exists?

The rest of this document answers that with a literature synthesis of the current bibliography (see §2), a map of select existing software, a reinventing-the-wheel test, and combinations that are complementary rather than redundant.

The lens for that synthesis is not “how to ship a weekend agent.” It is building a **local, multilingual (fr / en / pt), vendor-free, modest hardware compatible** tool covering heterogeneous ingest, indexing, curation, programming assistance, and exports (books, posts), all usable by both a human and an autonomous agent — **without closing doors** that a later year of work might need.



0. VERDICT IN ONE PAGE

0.1 WHAT THE LITERATURE ACTUALLY DEPICTS

Two claims, not one.

The **agentic-harness books** (Bhagwat, Berryman, Kar, Albada, Dibia, Labaschin, Lanham) still converge on a five-layer stack. That stack is the right picture for the original v4 question (hooks, tools, skills, MCP, DSL). Collapsing those five names into one word remains the error to refuse.

Layer	Books call it	What it is	What it is not
1. Model	LLM, completion, Responses API	A stochastic function: messages in, tokens out	The control plane
2. Prompt assembly	context engineering, application loop, packing	Code that <i>builds</i> the next request (history prune, retrieve, inject)	“Prompt engineering” as a chat craft


Layer	Books call it	What it is	What it is not
3. Tool surface	tools, function calling, actions, computer-use	Named operations the model may <i>request</i>	Authorization; the world
4. Transport	MCP, OpenAPI, RPC, SDK	How a host <i>discovers and invokes</i> a remote tool	Semantics; memory; the graph
5. Procedure pack	skills, playbooks, SKILL.md, CLAUDE.md	On-demand instructions + scripts, loaded when relevant	An operating system

The books do **not** treat MCP, skills, and tools as synonyms. That is still true after the fuller shelf.

The rest of the bibliography do not replace that stack. They show it is too small to be “what the literature depicts” for Projet Complexe. Prompt-pack manuals (Gautam, Sonvane) even depict a *competing* one-layer world — better wording — which Berryman already rejected. Keeping the five layers as the *whole* map would close the doors §2 exists to keep open: language, modest hardware, curated knowledge, exports.

What the expanded shelf actually adds (see §2 for file-by-file):

Plane	Who depicts it	What it is	What it is not	Door it keeps open
A. Model routing	Moslem & Kelleher 2603.04445; Ozdemir; Labaschin economics	Pick or cascade among <i>independently trained</i> models per query	One vendor identity; MoE inside a single net	Modest hardware <i>and</i> a later stronger model
B. Retrieval / index	Norman; Magda; Stewart; v3	Lexical + selected vectors + metadata over canonical files	Memory; the conceptual graph; “just packing”	Heterogeneous ingest without stuffing the window
C. Language system	Yu & Yao	Ingest → terms → generate → QA → versioned export; locale is a contract	“The model speaks French”	fr / en / pt as Requirements, including books and posts
D. Knowledge / fusion	Devlin (as recipe); Koch; Long 2603.24402; Mistrik; Kolade	Claims, typed links, gaps, uncertainty, HITL commit	GraphRAG summaries; memory MCP; a paper mill	The research loop without automating the degree
E. Learning / orientation	Dupoux et al. 2603.15381; Lanham confirm-vs-autonomous	Observe vs act vs <i>switch</i> (System M ≈ killswitch)	Chat logs as “the agent learned”	Later real learning; task vs research now



Read 0.1 as: **five layers for how agents get tools; five more planes for what this project is for.** The hook/DSL idea lives in the first slice. The ambition ("projet complexe", three languages, local, doors unclosed) lives in the second. Neither is a reason to merge MCP with skills, or RAG with knowledge.

0.2 WHAT THE IDEA GETS RIGHT

Wrapping LLM calls in **pre- and post-processing** is not a new agent architecture. It is the *definition* of a harness:

- Berryman & Ziegler: the product is an **application loop** that snippetizes, retrieves, and assembles the next prompt.
- Bhagwat: **tool design is the most important step**; middleware/guardrails sit *around* the loop, not inside the system prompt.
- Labaschin & Wallace / Grootendorst: **context engineering** is packing, not “memory.”
- Kar (2026): orchestration patterns (router, HITL, planner-executor, watchdog) are compositions of *named steps with I/O*, not a smarter model.
- Pi (@earendil-works/pi): transformContext → convertToLlm → provider; events before_provider_request / after_provider_response; registerTool; Agent Skills.

The fuller shelf does not change that mapping. It says *what else* those same hooks must be allowed to do, or the wrap is a toy:

- pre_llm **is also the router** (plane A, 2603.04445): which Technology, local-first, cascade if the small model fails. Not a second product.
- pre_llm **is also the language packer** (plane C, Yu & Yao): locale of the query, of the chunks, of the export contract. Not “the model speaks French.”

- `post_llm` is also the **commit gate** (plane D/E): propose a Claim / a `publish` draft; HITL or killswitch; never multi-agent consensus as truth (2603.24402, Kolade).

ASC already has the generic mechanism this maps onto: **hooks** (`pre_*` / `post_*` / variants), **entry points** (filename-safe, YAML-described), **DSL** (positional / boolean / named options that compile to bash). An LLM invoke is just another action. Tools the model may request are just other actions. That mapping is sound — and incomplete if hooks only “tweak the prompt.”

0.3 WHAT THE IDEA GETS WRONG IF TAKEN LITERALLY

Literal reading	Why it fails
“Replace MCP”	MCP is a <i>transport</i> (JSON-RPC + JSON Schema). You do not replace USB-C by inventing filenames. You decide whether you <i>need</i> that plug for third-party tools.
“DSL instead of JSON/Python”	Frontier models emit JSON tool calls . YAML <code>able.yml</code> + DSL is the <i>source of truth</i> for ASC. JSON Schema is a projection of that source for the model. Replacing the projection is fighting every provider SDK.
“Implement what agents’ tools/skills are”	Tools ≠ skills. A tool is an executable with a schema. A skill is a <i>procedure pack</i> (progressive disclosure). ASC already has both shapes (entry point vs documentation/hook recipe). Do not merge them.
“ASC becomes the coding agent”	That is how you accidentally rebuild Pi, Cursor, Claude Code, Goose, and Codex. ASC is glue. Coding-agent loops are a Technology behind <code>run-agent</code> .
“Hooks around the LLM are Projex Complexe”	The wrap is necessary and small (ASC). The product is ingest, curate, export, research — in three languages, on modest hardware. A perfect <code>pre_llm</code> with an English-only index has closed the door the shelf just opened.
“One model behind <code>llm</code> ”	Moslem: routing/cascading among independently trained models. Freezing Claude, GPT, or a single 7B as identity repeats Clinton’s obsolete-guide failure.

0.4 RECOMMENDATION (USABLE BEYOND ONE PKM PROJECT)

Do explore the idea — as a thin composition, not as a second MCP and not as a second Pi.

1. **ASC-generic (put in core, or keep as doctrine)**: every LLM call is an entry point wrapped by `pre_llm` / `post_llm` (and variants). Tool names the model sees are **allowlisted entry points**. YAML `entity/able` is canonical; JSON Schema / MCP tool lists are generated. DSL addresses arguments (`p1, o-max-4`). The model never sees `make hook`.
2. **Delegate (do not reimplement)**: multi-provider streaming (`pi-ai` or equivalent), coding-agent TUI and session trees (Pi / Cursor), Tree-sitter code graphs (code-graph-rag or similar), MCP as a *protocol*, Agent Skills `SKILL.md` as an *interchange format* if you want other harnesses to load the same procedures.
3. **Project Complexe ASC (instance-specific)**: which tools exist (`index`, `extract`, `relate`, `research`, `run-agent`, `publish`, ...); packing from v3 **plus locale**; the router/cascade policy (local 7B default, overflow allowed or not); killswitch; HITL accept of claims and of exports. These are compositions, not a new protocol.
4. **Refuse**: MCP as the user-visible capability surface; `CLAUDE.md` / skill folders as a second OS; unbounded tool catalogs; computer-use as default; Python/TS tool registries that bypass YAML; teaching the model a private DSL as its native function-call language; English-only embeddings as the store; a prompt pack as `publish`; AI-Supervisor-style paper mills; WhatsApp/OpenClaw as the control plane.

The rest of the note is the evidence for that cut.

1. THE QUESTION, RESTATED WITHOUT HYPE

1.1 WHAT WE ALREADY HAVE IN ASC

ASC’s job is to make computational things **nameable, addressable, composable, and executable**. For agents, that already implies:

- A **pivot** (`$subject-$action`) is a named capability. `run-agent`, `extract`, `relate` are pivots, not MCP tool names.
- An **entry point** is a file that can be invoked. Hooks wrap it (`pre_extract`, `post_extract`, `pre_extract/pdf`).
- **DSL** is a compact, filename-safe encoding of argv. It is for *addressing and validation*, not for chatting with a model.
- **YAML** (`*.entity.yml`, `able.yml`) describes what is required and how to validate. The DSL example in the README (`test-in(p1,[slug(p1),...])`) compiles to bash. That is already “schema as code,” just not JSON Schema.

So the proposal is not “add hooks.” It is: **treat the LLM itself as an entry point**, and treat **every tool the model may call as another entry point**, so the same pre/post machinery, the same allowlist, the same traces, apply to both.

That is a good genericity test. If it only works for Projex Complexe’s second brain, it does not belong in ASC. If it works for *any* project that calls a model and must not let the model invent `rm -rf`, it does.

1.2 WHAT “TOOLS / ACTIONS / SKILLS / MCP” NAMES IN THE WILD

These four words are used interchangeably on social media. The books and the working software do not.

Tool / function / action. A named operation with arguments, usually JSON Schema, executed by the *host* after the model requests it. OpenAI function calling, Anthropic `tool_use`, Google function declarations, LangChain `@tool`, Pi `registerTool`, MCP `tools/list` + `tools/call`. Bhagwat: designing this surface is the main product work.

Skill. A *reusable capability pack* loaded on demand: markdown instructions, optional scripts, optional pointer to tools. Anthropic Agent Skills (`SKILL.md`), Pi skills, Claude Code skills. Progressive disclosure: only the *description* sits in the system prompt; the body is read when needed. This is closer to an ASC **procedure** or a Projex Complexe **Requirement/Fallback** file than to a pivot.

MCP (Model Context Protocol). JSON-RPC 2.0 between a **host** (owns the loop, the model, the allowlist) and a **server** (exposes tools/resources, contains no model). USB-C metaphor: N+M instead of N×M. It does not decide *which* tools exist. It does not store knowledge. It does not replace tool-use training (ToolRL). Labaschin already warned not to bet the farm on it remaining the only norm.

Harness. Two meanings, both useful:

- **Eval harness:** fixtures, scores, ablations (`datasciencebrain`, RAGAS, Winteringham).
- **Agent harness:** the loop around the model (Pi, Cursor, Claude Code, Codex). Pi’s own name for itself.

ASC is already a harness for *the machine*. It is not yet an agent harness, and it should not become a *coding* harness. It should **host** one.

1.3 THE PRECISE CLAIM TO EVALUATE

It is theoretically possible to implement tools/actions/skills as ASC entry points, and to replace MCP with simpler ASC actions addressed in DSL rather than JSON/Python.

Split it:

Sub-claim	Verdict
LLM call as a hooked entry point (pre/post prompt)	Yes. Do this. It is the portable form of every harness.
Tools as allowlisted entry points	Yes. Do this. This is Bhagwat’s tool design, named in ASC.
Skills as on-demand procedure files	Yes, as files — not as a new runtime. Optionally <i>also</i> emit <code>SKILL.md</code> for Pi/Claude.
Replace MCP	No, as a protocol. Yes, as the default local glue. Local tools should be ASC. MCP is an adapter for <i>other people’s</i> tools.
DSL instead of JSON/Python	No at the model boundary. Yes as ASC addressing. Compile YAML+DSL → JSON Schema / MCP tool list / TypeBox. Never the reverse as source of truth.

This table answers **only the harness cut** (layers 1–5 in §0.1). It does not settle planes A–E (router, index, language, knowledge/fusion, orientation). Those are PCA/Projet Complexe policy. Do not read a “yes” on hooks as “projet complexe is done.”

2. LITERATURE REVIEW SUMMARY


This is not a second agents-literature-review. It is a **door inventory**: what each file makes newly thinkable for Projet Complexe, and which architectural shortcuts would lock a later use out.

2.0 AMBITION LENS (WHY “DO NOT CLOSE DOORS”)

"Projet Complexe" is a tool aiming to be useful for studying and for being more productive (with a beautiful and reactive desktop UX) : a place where **heterogeneous sources** (PDFs, notes, web captures, code, audio later) can be ingested and indexed; where **French, English, and Portuguese** are first-class for retrieval *and* for exports (a book chapter, a post, a bilingual note); where **curation** (accept / contradict / gap) is the knowledge act; where **programming assistance** is a worker, not the institution; where a **local agent** can use the same named actions a human uses from the terminal — on **modest hardware** (laptop, LAN box, 16 GB dedi), **without a vendor as the brain**.

Closing a door means making an identity choice that a later year cannot undo without a rewrite: English-only embeddings as the store; a cloud chat product as memory; RAG as knowledge; MCP as the vocabulary; one graph database as the conceptual model; computer-use as “agency”; a prompt pack as the product.

Keeping a door open means **naming the capability** (ingest, index, export, curate, code-assist, research, run-agent) and leaving the Implementation swappable. Language is metadata on sources and on exports, not a fork of the app.



Modest hardware is not a reason to close the frontier-model door. It is a reason to make **routing and cascading** first-class (paper 2603.04445): a local 7B for routine packing and research; escalate only when the Task demands it. Cloud remains a metered overflow, never the SoR.

2.1 CONSOLIDATED TABLE

Columns: **interesting** = what is not just another “build an agent” chapter. **For Proje^t Complexe** = how it bears on the ambition above (keep / adapt / do not let this become identity). Duplicate files of the same work are one row.

Source	Interesting / original	For Proje ^t Complexe
Stewart & Huang, <i>Agentic AI Data Architectures</i> (2026)	“Memory as infrastructure”: agents fail because <i>data</i> is fragmented (structured / unstructured / temporal), not because the model is small. Argues for distributed SQL as unified retrieval, semantic-transactional joins, episodic vs long-term stores, governance in the fetch path.	Steal the diagnosis. Adapt as Postgres (Magda, v3), not Cockroach-at-home. Keep the door to a later distributed SQL if the corpus outgrows one box. Refuse “the database is the agent’s mind.”
Norman, <i>Agentic RAG Systems</i> (2026)	Production RAG: why naive pipelines fail (~“40%”); lexical vs semantic vs structural similarity; chunking; hybrid fusion; rerank; GraphRAG (multi-hop ceiling, Microsoft community-summary pattern, reindex pain); Self-RAG / Corrective RAG; multi-agent retrievers; RAGAS + observability + FinOps.	Steal hybrid cascade and “static retrieval cannot recover from its own mistakes” → agentic <i>research</i> loop with a budget. Adapt GraphRAG as an Implementation of relate/research , not as the knowledge model (v2/v3). Refuse LangGraph as host; refuse community-summaries as Claims.
Lanham, <i>AI Agents in Action</i> (2025)	Spectrum: no agent → proxy → confirm then execute → autonomous. Function calling as “model proposes, runtime does.”	Steal confirmation as a state (HITL). Keep the door to later autonomy as an <i>explicit</i> Requirement, not the default. Refuse “the model asked” as authorization.
Gautam, <i>AI Prompt Engineering Bible</i> (2026)	Prompt-craft manual (how to write prompts that “work”).	Peripheral. Steal nothing as architecture. Keep a prompt library as files if useful; do not let it become the product.
Sonvane, <i>AI Prompt Mastery Guide</i> (2026)	850+ ready-made prompts (image, video, book covers, posts).	Folklore. Useful only as <i>examples of export genres</i> (cover, post, carousel). Refuse as a capability surface. Closing this door (not shipping a prompt shop) opens the door to real publish pipelines.
Osmani, <i>Beyond Vibe Coding</i> (2025)	Spectrum from vibe-coding to AI-assisted engineering; 70% problem; AI as drafter / pair / validator; when to stop vibing.	Steal: programming assistance is a mode with a human editor-in-chief. Keep the door to Pi/Cursor as workers. Refuse the coding agent as the second brain.
Osmani, <i>Vibe Coding</i> (2025, early release)	Same thesis, shorter. Overlaps <i>Beyond</i> .	Treat as the same door as the row above.

Source	Interesting / original	For Projet Complexe
Raschka, <i>Build a Reasoning Model (From Scratch)</i> (2026)	How reasoning models are <i>trained</i> (not how to prompt CoT). Weights, data, recipes.	Opens a late door: a local reasoning small model as Technology. Do not close it by assuming only API “o-series” can research. Do not open it in v1 (training is not an immediate “projet complexe” goal).
Ozdemir, <i>Building Agentic AI Workflows</i> (2025)	Workflows vs agents; evals; multimodality; reasoning models vs computer-use ; fine-tune; compression; Matryoshka embeddings; local-ish case studies (Qwen speculative decode, Moondream VQA).	Steal: explicit tools > computer-use; multimodel; compression for modest hardware. Adapt computer-use as a sandboxed later Implementation. Keep the door to image extract (v3 figure profile) without making AIGC a default tool.
Albada, <i>Building Applications with AI Agents</i> (2025)	Holistic lifecycle: when to use an agent; tool design; memory; single vs multi-agent; learning from experience; evals; security. The book the author “wanted to hand colleagues.”	Steal the <i>questions</i> (good tools, when RAG, when multi-agent). Adapt answers as ASC pivots. Refuse another framework ontology.
Devlin, <i>Building LLM Agents + ...with RAG, KG, and Reflection</i> (2025)	Two files, one project: RAG + graph + reflection as the agent recipe.	Steal reflection as <code>post_llm</code> / a second hooked call. Adapt KG queries as <code>relate</code> on <i>accepted</i> links. Refuse RAG+KG as the knowledge plane (v3).
Kar, <i>Building Multimodal Generative AI and Agentic Applications</i> (2026)	Catalog of RAG, vectors, guardrails, agents, MCP , HITL and a long pattern list (router, watchdog, planner-executor, database-with-tools...). Local Ollama chapter.	Steal guardrails-as-layer, HITL, “tools talk to the DB.” Adapt each pattern as a <i>composition</i> , not ontology. MCP = one chapter, not the architecture. Keep local-Ollama door.
Clawdbot / OpenClaw handbook (OCR, 2026)	Local-first agent on Mac Mini / Docker; WhatsApp/Telegram bridges; gateway websocket; prompt-injection hardening; “always-on” sizing.	Steal local-first + sandbox. Adapt “always-on” to the LAN box / dedi , not to a chat app. Refuse WhatsApp as control plane; refuse pairing-codes-as-identity.
Kolb & Rosen, <i>Cognitive Kin</i> (2026)	Working <i>with</i> agents; meaning stays human.	Steal collaborator model. Keep the door to “the human still judges.” Refuse personality-as-architecture.
Nolan & Stoudt, <i>Communicating with Data</i> (2021)	Writing for data science: how claims are <i>shown</i> , not retrieved.	Steal: export quality (a post, a chapter) is a writing craft. <code>publish</code> is not “the LLM dumped markdown.” Keep the door to human-edited books.
Sanderson et al., <i>Data Contracts</i> (2025)	Schema + owner + SLO + changelog. Silent drift is the outage. Shift-left quality.	Steal: YAML <code>able</code> is the contract for tools <i>and</i> for export schemas (a “post” and a “book chapter” are contracts). Keep the door to versioned exports.
Dibia, <i>Designing Multi-Agent Systems</i> (2025)	Taxonomy: workflow (explicit control) vs autonomous (emergent) ; UX for multi-agent; structured output; middleware; OTel; agents-as-tools; computer-use chapter; “but what about frameworks?”; eval harness.	Steal: default to explicit workflows (named ASC steps). Keep the door to later autonomous orchestration as an opt-in. Refuse swarms and computer-use as v1. UX principles → <code>inspect-agent</code> / killswitch, not a second web app inside the agent.
Torralba & Isola, <i>Foundations of Computer Vision</i> (2024)	Textbook: images, CNNs, transformers, vision+language.	Background. Keep the figure extract door (v3). Do not put a vision stack in the core ontology.
Brikman, <i>Fundamentals of DevOps</i> (2025)	Glue is a profession: scripts → IaC → orchestration. Combine tools.	Steal: ASC is this layer. Keep the door to Compose on modest hardware; do not require k8s for the laptop.
Kolade & Egbetokun, <i>Generative AI in Research</i> (2026)	LLMs in research design , analysis, feedback — the human/supervisor use.	Steal: the app is a research instrument , not a chat. Keep HITL on claims. Refuse unsupervised “the agent wrote my MA.”
Bartlett, <i>How to Talk to AI</i> (2026)	Jailbreaks, narrative entanglement, habits.	Steal: prompt as attack surface. Adapt as hooks, not pep talks.
Sayed, <i>Inference and Learning from Data</i> vols 1–3 (2023)	Three-volume mathematical core (foundations, inference, learning).	Background. Keep the door to <i>real</i> stats later (eval, uncertainty on links). Do not pretend RAG scores are Bayesian fusion (see Koch).
Yu & Yao, <i>Intelligent Language Services</i> (2026)	Most original book on this shelf for Projet Complexe. Language work is a system : ingest → terminology → generate → evaluate → revise. Text becomes a versioned, traceable knowledge asset. Translation agents, prompt-as-NL-programming, RAG for citation, when to LoRA, multilingual workflow design, quality contracts, ethics. Gravity moves <i>upstream</i> of “translate this.”	See §2.2. This is the fr/en/pt export and ingest doctrine. Steal the system. Do not become a CAT tool vendor.

Source	Interesting / original	For Projeet Complexe
Röber, <i>Interpretable ML</i> (2026)	Optimization-based explanations; human-centered eval.	Steal: <code>inspect-agent</code> and Factor displays. Keep the door to <i>why this chunk</i> .
Reddi, <i>Introduction to ML Systems</i> (2025)	Serving, batching, accelerators.	Steal: model as a service with SLOs. Keep Ollama/llama.cpp/API as Environments.
Magda, <i>Just Use Postgres</i> (2025)	One database: relational, JSON, FTS, pgvector.	Already v3 SoR. Stewart’s “distributed SQL” is the <i>scale door</i> this keeps closed until needed — correctly, on modest hardware.
Mistrik & Galster, <i>Knowledge Management in Data-Intensive Systems</i> (2023)	KM in engineering orgs: traceability, process, docs — not RAG.	Steal: knowledge is curated artifacts + process . Keep skills as files.
Labaschin & Wallace, <i>Managing Memory for AI Agents</i> (2025)	Memory types; multimodel economics; lock-in; MCP as more tools <i>and</i> as memory servers.	Steal cost-of-tools and “do not bet the farm on one protocol.” Refuse memory MCP as store.
Miller, <i>Mathematics of Optimization</i> (2017)	How to do things faster (classical optimization).	Background. Routing/cascading (paper 04445) is the applied form.
Sadhu & Konar, <i>Multi-Agent Coordination</i> (2021)	MARL, correlated equilibrium. Pre-LLM.	Steal: coordinate through the world (files, events). Refuse chat-bus swarms.
Lin & Liu, <i>Multimodal Large Models</i> (2026)	Architectures, VQA, AIGC, embodied, world models, open-source platforms.	Keep read multimodal (extract figures, later ASR). Close embodied / generate-video as identity. World-model talk → see paper 24402, not a robot.
Bhagwat, <i>Principles of Building AI Agents</i> (2024)	Tool design first; MCP as remote execution; workflow graphs when free-tool agents are wild.	Steal allowlist, traces, “compile the Task into steps.” Workflow DAG ≠ conceptual graph.
Bhagwat & Gienow, <i>Patterns</i> (2026)	Lethal trifecta; middleware guardrails; evals.	Steal perimeter vs inner loop. Dual authz. Sandbox.
Baihan Lin, <i>Privacy and Security for LLMs</i> (2026)	Metrics, membership inference, RAG privacy, DP, federated, secure deploy.	Steal: retrieved chunks are data flows; <code>pre_llm</code> redacts before any cloud hop. Keep local-default . Exotic crypto is a closed door for v1 (good).
Berryman & Ziegler, <i>Prompt Engineering for LLMs</i> (2025)	Prompting = application loop : snippetize, score, assemble, parse, tools.	Steal: <code>pre_llm</code> / <code>post_llm</code> are this loop. The hook idea is this book, named in ASC.
Innanen, <i>Prompted</i> (2025)	How to create and communicate with AI (practitioner/design).	Soft. Steal: communication is design. Refuse as a protocol.
Kofler, <i>Scripting Automation</i> (2024)	Bash / Python / glue; do one thing well.	Steal: entry points may be bash (CLI = GUI). Not the model-facing language.
Winteringham, <i>Software Testing with Generative AI</i> (2024)	Skeptical LLM use; contracts; TDD with models; test data.	Steal: eval harness around traces and <code>publish</code> outputs (fr/en/pt). Refuse LLM-as-only-judge.
Wong, <i>The AI Cybersecurity Handbook</i> (2026)	Security handbook for AI systems.	Steal: injection, supply chain, least privilege. Same family as lethal trifecta. Keep tools few.
Clinton, <i>The Complete Obsolete Guide to Generative AI</i> (2024)	2024 snapshot; title admits decay.	Meta-lesson: do not freeze a vendor stack as identity . The guide becoming “obsolete” is why ASC names capabilities.
Osmani, <i>The Effective Software Engineer</i> (2026)	NIH, tool obsession, over-engineering, hero complex.	The filter for \$3–\$5: wrap Pi, do not rebuild it.
Koch & Schlangen, <i>The Future of Information Fusion</i> (2025)	From Bayes to cognitive fusion : aleatoric vs epistemic uncertainty; trustworthiness of fused statements; not LLM-native.	See \$2.3 . The conceptual graph is closer to <i>fusion with uncertainty</i> than to cosine RAG. Keep the door to typed Links with Factors.
Paper 2603.04445 Moslem & Kelleher	Survey of routing and cascading across <i>independently trained</i> LLMs (not MoE). When / what / how of the decision. Can beat a single flagship on cost×quality.	See \$2.4 . This is how modest hardware stays compatible with later stronger models.


Source	Interesting / original	For Projet Complexe
Paper 2603.15381 Dupoux, LeCun, Malik	Why current AI does not <i>learn</i> like organisms. Systems A (observation), B (action), M (meta-control switching). Data wall; language-centrism.	See §2.5. Killswitch ≈ System M. Do not close the door of later real learning; do not fake it with chat logs.
Paper 2603.24402 Long, AI-Supervisor	Persistent Research World Model (KG with uncertainty on edges); gap discovery; consensus before commit; curiosity-driven supervision for people <i>outside</i> elite labs. Explicitly the human/supervisor problem.	See §2.5. Steal the <i>shape</i> . Refuse paper-mill autonomy and “GPU validation = truth.” HITL remains the commit.

2.2 LANGUAGE AS A SYSTEM (FR / EN / PT) — YU & YAO

Intelligent Language Services is the book that is actually about **exports and ingest in several languages**, which most agent books treat as an afterthought (“the model speaks many languages”).

Yu & Yao’s claim, restated: a translation or a blog post is the **surface** of a pipeline (ingest, segment, terminology, generate, evaluate, revise). High-value work sits *before* the sentence: corpora, terms, prompts-as-contracts, where only a human may sign. A text is not a terminus; it is a **versioned, searchable asset**. Models hallucinate → retrieval and governance compensate. Drift → workflow + structured output + quality contracts. Accountability → observability + human intervention.

That maps onto Projet Complexe without becoming a CAT suite:



Steal. Language services = system architecture. Terminology is a curated stock (Projet Complexe entities), not an embedding cluster. Multilingual export is a **pivot family** (publish with a language + genre contract), not a prompt from Sonvane’s chapter 3.

Adapt. fr/en/pt are the *required* locales for evals (Winteringham): a retrieval that only works in English has failed, even if RAGAS looks fine. Meilisearch and `tsvector` both care about language analyzers — that is Implementation detail, named as a Requirement. Translation agents (Yu & Yao ch. on task chains: ingest → register → terms → generate → multi-model) are **compositions of extract / relate / research / publish**, not a new product.

Keep open. Later: LoRA on a local model for a personal termbase; speech; a fourth locale. **Do not close** by picking one English embedding model as identity, or by storing only “the English chunk.”

Refuse. Fine-tuning as v1 (Raschka/Yu show it is a *project*, not a toggle). Becoming SDL Trados with an LLM. Per-word billing ontology.

2.3 RETRIEVAL, KNOWLEDGE, AND FUSION (NORMAN, MAGDA, STEWART, KOCH, DEVLIN)

Three different “graphs” keep getting one name:

Graph	What it is	Engine (v3 default)	Failure if confused
Lexical / vector index	Projection for packing a small window	Meilisearch + selected pgvector	Treated as memory
Conceptual / evidentiary	Claims, typed Links, unknowns, gaps	Postgres JSONB / link tables	Replaced by cosine or by GraphRAG summaries
Code structure	AST, calls, traces	Opt-in (code-graph-rag)	Used as the personal wiki

Norman is useful because he *admits* the multi-hop ceiling of text RAG and then sells GraphRAG. The steal is the ceiling; the refuse is Microsoft-style community summaries as the human’s knowledge. Devlin’s reflection is `post_llm`, not a third database.


Stewart’s “memory as infrastructure” is right as a complaint (stateless model + scattered files) and wrong as a prescription if it means a new agent-native distributed SQL. Magda already opened the **one-Postgres** door; Stewart keeps a later scale door; v3 stays on Magda until it hurts.

Koch & Schlangen are the unexpected ally: **information fusion** distinguishes noise in the data (aleatoric) from ignorance in the model (epistemic), and cares whether a fused statement is *trustworthy*. A Factor on a Link, a KnowledgeGap, a `valid_at` are closer to this than to an embedding. Keep that door: the graph can grow toward fusion without becoming a defence-lab Bayesian engine in v1.

2.4 MODEST HARDWARE: ROUTING, CASCADING, LOCAL MODELS

Moslem & Kelleher (`2603.04445`) is the paper that makes the hardware constraint *compatible* with not closing the quality door.

- **Routing:** one decision, pick a model (difficulty, domain, cost, cluster, RL, uncertainty).
- **Cascading:** try small/fast; escalate if quality/confidence is insufficient.
- Practical systems **compose** both. They can outperform a single flagship on the joint objective.



Ozdemir (compression, speculative decode, Matryoshka) and Kar (Ollama chapter) are the engineering companions. Labaschin’s multimodel economics is the same idea in product language. Raschka keeps a **train your own reasoner** door shut for now, open later.

Do not close: a future local 32B on the LAN box; a Portuguese-strong model as a named Technology; CPU-only llama.cpp on the laptop. **Do not open as identity:** “we only use Claude.” Clinton’s *Obsolete Guide* is the cautionary tale of identity-as-vendor.

2.5 THE HUMAN-RESEARCHER LOOP (KOLADE, COGNITIVE KIN, PAPER 24402, PAPER 15381)

This is the emotional core of the ambition, and the place the literature is most original *and* most dangerous.

Kolade: generative AI in *research design and feedback* — the human use. Steal as orientation: the desktop is a research instrument. Refuse: the agent as author of the degree.

Cognitive Kin: meaning-making stays human. The senior developer with two MAs is not trying to automate scholarship; they are trying not to lose it to ChatGPT’s amnesia.

AI-Supervisor (2603.24402): the paper that says out loud that research supervision is rationed by institutions, and that a **persistent world model of the literature** (methods, benchmarks, gaps, uncertainty-tagged edges, consensus before commit) would let a curious individual work without a funded lab. That is “the tool I wished I had as a student,” at research scale.

Steal the **shape**: a graph that survives sessions; gaps as first-class; corroboration before commit; curiosity as input. The paper’s own “lightweight exploration vs full-scale investigation” matches modest hardware.

Refuse: autonomous paper writing; “agents verify on GPUs therefore it is science”; replacing HITL with multi-agent consensus; treating OpenReview ingestion as a default network tool (lethal trifecta). Projet Complexe’s graph is **personal and curated**, not a publication mill.

Dupoux, LeCun, Malik (2603.15381): today’s models do not autonomously learn. They propose System A (observe), System B (act), System M (switch). That is already the killswitch / research-vs-task split, stated as cognitive science. Keep the door to *later* real learning (new observations, not more PDF chunks). Do not close it by stuffing chat into a memory product and calling it System A.

2.6 WHAT THE SHELVES AGREE ON, AND WHAT THEY LEAVE SILENT

Agree (do not reopen):

1. The model does not execute; a host does (Lanham, Bhagwat, Albada, Dibia).
2. Tools are few, named, schema’d, permissioned (Bhagwat, Albada, Winteringham, Wong).

- 3. Prompt assembly is code (Berryman). MCP is a plug (Kar, Labaschin), not a brain.
- 4. Naïve RAG fails; hybrid + metadata + evals are the floor (Norman, Devlin, v3).
- 5. Local-first and privacy are design, not a slogan (Lin, OpenClaw, Bartlett).
- 6. More tools raise capability *and* cost, injection surface, and eval burden.

Silent — where **Projet Complexe** / **ASC** can stay original without being cute:

- **Three working languages as a Requirement**, not a model marketing bullet (Yu & Yao get closest).
- Filename-safe **addressing** of tools and hook variants (ASC DSL).
- One vocabulary for **machines, processes, and model calls**.
- **Genericity**: promote a local name; do not start from an MCP registry.
- **Two (or three) graphs** that must not share a word (knowledge ≠ RAG ≠ code AST).
- A modest-scale Research World Model that **commits only with a human** (24402’s shape, Kolade’s ethics).

Disagreement: MCP enthusiasm (Kar, older Grootendorst review) vs lock-in caution (Labaschin) vs “don’t roll a client this year” (Bhagwat). **Resolution unchanged**: generate MCP if a neighbor host needs it; do not live there.

Prompt-pack books (Gautam, Sonvane, partly Innanen) disagree with Berryman: they sell wording; he sells an application loop. Closing the “850 prompts” door is how you keep **publish** open as a real pipeline.

3. HOW THE INDUSTRY IMPLEMENTED THE SAME IDEA (2025-2026)

3.1 THE HARNESS PATTERN (THIS IS PRE/POST ON THE PROMPT)

Pi documents the loop explicitly:

```
AgentMessage[] → transformContext() → AgentMessage[] → convertToLlm() → Message[] → LLM
                  (optional)                               (required)
```

transformContext: prune, inject. **convertToLlm**: drop UI-only messages, fit the provider.

Extensions subscribe to a lifecycle that is already the hook list ASC would invent:

Pi event (coding agent)	ASC analogue
before_agent_start (inject message, edit system prompt)	pre_run-agent / pre_llm
before_provider_request (inspect or replace payload)	pre_llm variant per provider
after_provider_response	post_llm
tool_call (block or modify)	pre_<tool-entry-point> — authz lives here
session_before_compact	pre_compact (already named in the memory review)
registerTool + TextBox schema	YAML able → generated JSON Schema
Agent Skills SKILL.md	procedure files / Requirements; optional export

Cursor, Claude Code, Codex, Goose, OpenHands, Aider all have some subset of: system prompt files, tools, permissions, compaction, MCP or native tools, skills.

Implication. Exploring “hooks around LLM calls” is **not** original as a *behavior*. It is original as a **language-agnostic, filename-addressable, YAML-canonical** form of that behavior. That originality is real and small. It is ASC’s actual job. Rebuilding the TypeScript agent runtime, the TUI, the session tree, and the provider catalog is **not** ASC’s job. That is Pi (and Cursor, etc.).

3.2 MCP (THE PROTOCOL YOU WOULD BE “REPLACING”)

As used in the agents literature review and the social-media review:

- Server: tools + JSON Schema, no model.
- Host: loop, model, aggregation, **allowlist** (if the host is serious).
- Client: one connection per server.

What MCP solved: N models × M tools glue; language-independent servers; a way for Cursor/Claude to attach GitHub, Postgres, browsers.

What MCP did not solve: which tools *should* exist; authorization; prompt injection via tool output; memory; knowledge; evals. Bhagwat: it is remote code execution with a nicer brochure.

code-graph-rag ships an MCP server so Claude Code can query the code graph. That is the correct use: a **specialized Implementation** speaking a plug the *neighbor host* already has. If ASC is the host, that project’s **CLI/SDK** (**cgr**) is enough; MCP is for when Cursor is the host.

3.3 AGENT SKILLS (THE OTHER STANDARD)

Pi implements the **Agent Skills** convention (**SKILL.md** + frontmatter **name** / **description**, progressive disclosure, optional **allowed-tools**). Claude Code and Codex have parallel folders. Pi even loads **~/.claude/skills** on purpose.

Skills are **not** MCP. A skill is a brief. It may *tell* the model to call tools or to run a script.

Implication. If the goal is interoperability with other harnesses, **export procedure files as SKILL.md**, do not invent a third packaging format. If the goal is execution under ASC, **keep the executable as an entry point** and the brief as markdown/YAML. Two objects.

3.4 PI (EARENDIL-WORKS/PI) — WHAT TO STEAL VS WRAP

Pi is a **minimal terminal coding harness**: @earendil-works/pi-ai (providers), pi-agent-core (loop, tools, state), pi-coding-agent (CLI), pi-tui, telemetry. Extensions in TypeScript. Skills. Prompt templates. Packages. RPC/JSONL. Explicitly **no** built-in permission sandbox (containerize: Docker, Gondolin, OpenShell).

Pi owns	ASC should
Multi-provider streaming, model catalog	Call Pi or Ollama as a Technology; do not fork pi-ai
Coding tools (read/edit/bash)	Allowlist <i>which</i> of those exist when Pi is the Implementation of run-agent in task/code mode
Session JSONL, compaction, TUI	Stay Pi’s problem; import summaries as Notes if needed
transformContext / provider hooks	Mirror the idea as ASC pre_llm / post_llm so non-Pi models (a 7B on Ollama for research) get the same packing
TypeScript extensions	Contrib <i>or</i> an adapter: an extension that turns ASC entry points into registerTool
Agent Skills	Consume/emit; do not make ASC a skill runner that reimplements Pi

Pros of delegating to Pi: MIT, small core, self-extensible, already the hook/tool/skill split, RPC for embedding, huge mindshare, llama.cpp path.
Cons / costs: Node/Bun toolchain; no native sandbox; TypeScript-centric extensions; coding-agent default ontology (files, diffs) is **not** Projet Complexe’s Claim/Link ontology; permission model is “you are the user.”
Rule. Pi is a **worker**. ASC remains the host that *may start* Pi, pass a packed prompt / skill subset / allowlist, and receive events. Projet Complexe never talks to Pi’s RPC from the webview.

3.5 CODE-GRAPH-RAG (VITALI87) — WHAT TO STEAL VS WRAP

code-graph-rag parses a **source monorepo** with Tree-sitter (and ast-grep for some languages), stores a **code** knowledge graph in **Memgraph**, optional vectors in **Qdrant**, NL→Cypher, AST edit, dead-code, runtime **CALLS** via tracing, MCP server. This is a different graph from Projet Complexe’s **conceptual** graph (v2/v3). Revival v2 already parked it as a **later inner zoom for task-mode**, not day one.

It owns	ASC / PCA should
Multi-language AST, call graphs, runtime traces	Not reimplement
Memgraph + Docker daemon	Keep as <i>its</i> Compose; do not make Memgraph the Projet Complexe SoR (v3: Postgres)
NL query over code	Implementation of a pivot like relate / research when the object is a git tree
MCP server	For Cursor/Claude as neighbor hosts. For ASC-as-host, prefer cgr CLI/SDK behind an entry point
Surgical AST patch	Task-mode only, HITL diff, never knowledge-mode writes

Pros of delegating: years of parser work; mixed-language monorepos; runtime overlay; MCP already done.
Cons: heavy (Memgraph, Qdrant, Docker) — **opt-in only on modest hardware**; code-only; Cypher-from-NL is another stochastic layer (evals!); enterprise upsell; a second graph database next to Postgres invites the collapse v3 just refused. Default-on would close the laptop door \$2.4 is trying to keep open.
Rule. Use it when the Task is “understand or edit *this codebase*.” Do not use it for PDFs, claims, or Wikipedia. Do not dual-write accepted Claims into Memgraph.

3.6 OTHER PROJECTS IN THE SAME NEIGHBORHOOD (“ETC.”)

Project	Role	Overlap with the idea	Delegate?
Cursor / Claude Code / Codex	Full coding hosts + MCP + skills + hooks (PreCompact, etc.)	They <i>are</i> the productized version of “harness + tools + skills”	Yes, as neighbor providers . Never as the second brain.
Goose (Block)	Local-first agent, recipes, extensions	Same harness family	Optional alternative Implementation of run-agent
Aider / OpenHands	Repo-oriented agents	Coding loop	Same: Technology, not ASC
LangChain / LangGraph / LlamaIndex / Mastra	Frameworks: tools as Python/TS functions, graphs as code	They want to <i>be</i> the host	Spike a pivot, then throw away the framework as control plane (already the literature-review stance)
Composio / Pipedream / agentic iPaaS	SaaS tool glue	MCP/OpenAPI catalogs	Refuse as architecture (Bhagwat). Maybe one scoped adapter later
Letta / Mem0 / Zep / agentmemory	Memory products, often MCP	Compete with typed artifacts	Refuse as control plane (already decided)
qmd (tobi)	CLI hybrid search + MCP	Lexical+vector packer	Interesting <i>behind</i> research; v3 already chose Meilisearch+pgvector
Docling / GROBID	Extract Implementations	Tools the <i>extract</i> pivot calls	Already v3

Project	Role	Overlap with the idea	Delegate?
Ollama / llama.cpp	Local model serving	One Technology behind the router , not “the” model	Yes
OpenClaw / similar always-on chat agents	Local-first + messengers	Overlaps local-first; competes as control plane	Steal sandbox; refuse WhatsApp as host (see §2 table)
CAT / translation suites	Multilingual workflow products	Overlaps Yu & Yao’s system	Steal pipeline shape; refuse becoming Trados

None of these replace **YAML-canonical, filename-addressable tools**. All of them replace a **from-scratch agent loop** if you let them.

4. WHEEL-REINVENTING TEST

Osmani’s NIH anti-pattern, applied brutally.

4.1 WHAT WOULD BE REINVENTING

If you build...	You are competing with...	Verdict
A new JSON-RPC tool protocol	MCP	Refuse. Adapter only.
A TypeScript coding agent with TUI, sessions, compaction, providers	Pi	Refuse as ASC core. Wrap.
Tree-sitter → graph DB → NL Cypher for code	code-graph-rag	Refuse as default stack. Wrap when the object is code.
Python <code>@tool</code> registry as source of truth	LangChain, Mastra, Pi TypeBox	Refuse. <code>YAML_able</code> is source of truth.
<code>SKILL.md</code> runner + discovery + frontmatter	Agent Skills / Pi / Claude	Do not fork. Emit/consume the format.
Memory MCP / wiki daemon	Letta, Mem0, Karpathy gist forks	Already refused in the literature review.
Computer-use / Playwright from the UI	Ozdemir’s skip-tools path; lethal trifecta	Refuse as default.
A CAT / translation product	Trados-class; Yu & Yao’s industry	Refuse as identity. Steal the <i>pipeline</i> ; <code>publish</code> stays a pivot.
An autonomous “AI scientist” / paper mill	AI-Supervisor (<code>2603.24402</code>) as product	Refuse autonomy. Steal persistent graph <i>shape</i> + HITL commit.
Agent-native distributed SQL	Stewart & Huang	Refuse on modest hardware. Postgres (Magda) until it hurts.
A prompt shop / prompt IDE	Gautam, Sonvane	Refuse as capability surface.
A standalone model-router product	RouteLLM-class; paper <code>2603.04445</code>	Refuse as a new app. Policy lives in <code>pre_llm</code> / PCA YAML.

4.2 WHAT IS **NOT** REINVENTING (THE ACTUAL ASC-SHAPED GAP)

The ecosystem’s tools are **language-homed** (Python decorator, TypeScript `registerTool`, MCP server in any language) and **host-homed** (Cursor’s allowlist, Pi’s extensions). There is still no widely used layer that says:

The same naming system that starts a container, validates a YAML entity, and runs a hook also **names the model call and the tools the model may request**, in a filename-safe DSL, with pre/post hooks, independent of whether the model is 7B local or a frontier API.

That layer is **small**. It is:

1. An entry point `llm` (or `complete`, or `run-agent`’s inner call) with `pre_llm` / `post_llm`.
2. A generator: `able.yml` → JSON Schema (and optionally MCP `tools/list`).
3. A dispatcher: model tool name → allowlisted entry point → hooks → result back into the next prompt (another `pre_llm`).
4. Traces that name pivot, Technology, Environment, Task id — not spans dumped into the conceptual graph.

If that is all, it is **glue**, which is ASC. If it grows a TUI, a skill marketplace, a Memgraph, or a prompt IDE, it has become a product that already exists. That is the **ASC-shaped** gap. It is not the whole Projet Complexe gap. Language contracts, locale-aware indexes, `publish` genres, Factor-on-Link fusion, and the human HITL graph are **PCA / Projet Complexe** work. They are also not Pi, not MCP, and not a reason to fatten ASC. Split the two gaps or you will rebuild a coding harness *and* a CAT tool in the same repo.

4.3 “DSL INSTEAD OF JSON” — THE PRECISE SPLIT

ASC DSL is good at **argv**:

```
p1, p2, a
b-online, bo-y
o-max-4
```

JSON Schema is good at **nested records** (the shape models are trained to emit).
Compilation, not replacement:

```
able.yml + DSL validate → JSON Schema → provider tool list
                        → MCP tools/list  (optional adapter)
                        → TypeBox          (if Pi is the worker)
                        → filename / hook variant
```

Teaching a 7B model to emit `o-max-4` as a function call is a **fine-tune / constrained-decode research project**, not a revival decision. Even then, the canonical description remains YAML.
Exception (small): for *tiny* local models with no native tool calling, `post_llm` can parse a **constrained** mini-format (even DSL-like) *as a Fallback Implementation* of “the model requested a tool.” Do not make that the protocol you show Claude.

5. COMBINATIONS THAT ARE COMPLEMENTARY (NOT MUTUALLY EXCLUSIVE)

Mutual exclusion is rarer than people think. Redundancy is the real waste.

5.1 COMPLEMENTARY (DO COMBINE)

Combination	Why it is not redundant
ASC hooks + Pi	ASC packs/allowlists/traces; Pi streams, edits, sessions. Different jobs.
ASC entry points + MCP adapter	Local tools are files. <i>Foreign</i> tools that only speak MCP get a bridge. One catalog in YAML.
ASC procedures + Agent Skills export	Same brief, two skins: ASC docs vs <code>SKILL.md</code> for Pi/Claude.
v3 retrieval + <code>pre_llm</code>	Meilisearch/pgvector/Postgres are <i>what</i> you pack. Hooks are <i>when</i> .
code-graph-rag + <code>relate</code> on a repo	Code graph is an Implementation. Conceptual graph stays Postgres.
Lanham confirm + Bhagwat lethal trifecta	HITL on write tools; no private corpus + untrusted web + outbound in one catalog.
Data contracts + Winteringham evals	Schema on tools; tests on traces.
Ollama <code>research</code> + Pi <code>run-agent (code)</code>	Small local model for knowledge-oriented packing; frontier/coding harness for acting on a repo. Multimodel as Grootendorst said. This is already routing (<code>2603.04445</code>), even if the “router” is a named pivot rather than a learned classifier.
v3 retrieval + <code>locale filters</code> + <code>publish</code> contracts	Indexing is not exporting. Yu & Yao: same system, different stages. Nolan: the chapter/post is a writing artifact.
HITL commit + Koch-style Factors	Consensus of agents (<code>2603.24402</code>) is not a substitute. Uncertainty on a Link can wait; the <i>slot</i> for it should exist.

5.2 REDUNDANT (PICK ONE PER ROLE)

Pick one	Do not also...
YAML <code>able</code> as tool source of truth	Hand-written MCP server schemas <i>and</i> LangChain tools <i>and</i> Pi TypeBox as independent sources
Pi or Cursor or Claude Code as the <i>coding</i> worker for a given Task	All three in one loop (handoff via files/Notes, not via shared memory MCP)
Meilisearch lexical (v3)	Solr <i>and</i> always-on Postgres FTS <i>and</i> qmd as three live lexical brains
Postgres relations (v3)	Memgraph for <i>claims</i> (Memgraph only for <i>code</i> if code-graph-rag is on)
ASC as host	Mastra/LangGraph as host
Skills as on-demand briefs	<code>CLAUDE.md</code> as a second constitution
Postgres as SoR (v3)	Stewart-style distributed SQL <i>and</i> a memory MCP <i>and</i> Letta as three minds

Pick one	Do not also...
Conceptual graph (Claims / Links)	GraphRAG community summaries as the wiki

5.3 MUTUALLY EXCLUSIVE (ACTUALLY)

A	B	Why
Model never sees <code>make hook</code>	Renderer invokes arbitrary hooks	Lethal. Already refused.
Research orientation (no act)	Catalog includes <code>bash</code> , mail-send, computer-use	Killswitch becomes theatre.
Extract-once canonical files	Memory daemon rewriting notes from every tool call	Literature review already refused.
JSON Schema as <i>projection</i>	DSL as the provider’s function-call language for frontier models	Fights the ecosystem; split the Fallback for tiny local models only.
Canonical text keeps source language	Store only an English embedding / English chunk	Closes fr/en/pt (Yu & Yao). Locale is metadata, not a write-down.
HITL commit of Claims / exports	Multi-agent consensus as the writer	Kolade + <code>2603.24402</code> refuse: that automates the degree.
Local files as SoR	Cloud chat / OpenClaw messenger as memory	Lin, OpenClaw-as-host refuse.


6. WHERE THIS LIVES: ASC VS PROJET COMPLEXE ASC VS REFUSE

Genericity test from v2: would another project, not a second brain, need this?

Piece	Home	Why
Hook wrap around <i>any</i> named action, including <code>llm</code>	ASC core (doctrine now; code when an Implementation needs it)	Generic. Useful for deploys, not only agents.
DSL as argv addressing	ASC core (exists)	Do not extend DSL into a JSON competitor.
YAML → JSON Schema generator for tools	ASC core or a tiny generic helper	Any agentic ASC project needs it.
MCP <i>client</i> adapter (optional)	ASC contrib or PCA if only Projet Complexe needs GitHub	Generic but not v1-critical.
MCP <i>server</i> exposing ASC entry points	Later, read-only, scoped	Dangerous; useful if a neighbor host must call you.
Allowlist of tools for <code>research</code> vs <code>run-agent</code>	Projet Complexe ASC	Instance policy.
Packing recipe (Meilisearch k, token budget, Flow)	Projet Complexe ASC	Uses v3 stack.
Locale on extract / index / pack / <code>publish</code>	Projet Complexe ASC + Projet Complexe objects	Yu & Yao. Not generic ASC; every PCA instance may choose locales. fr/en/pt are <i>this</i> instance’s Requirement.
Router / cascade policy (which Technologies, when to overflow)	Projet Complexe ASC	Paper <code>2603.04445</code> as policy, not a new daemon.
Killswitch, Claim parsing in <code>post_llm</code> , HITL commit	Projet Complexe ASC	Semantic; System M (<code>2603.15381</code>).
Terminology / accepted entities	Projet Complexe	Interpretation. Agents propose; humans accept.
<code>publish</code> genre contracts (note, post, chapter)	Projet Complexe ASC	Data-contract shape (Sanderson) on exports, not Sonvane prompts.
Pi / Cursor / Ollama as Technologies	PCA environment YAML	Swappable Implementations.
code-graph-rag Compose	PCA, opt-in profile	Heavy; code-only.
Skill marketplace, MCP registry in the UI	Refuse	Tool obsession; lethal trifecta.
ASC as MCP host <i>visible to the webview</i>	Refuse	Thin Tauri.

7. A CONCRETE (STILL NON-SPEC) SHAPE

This is a picture to argue with, not an implementation plan.



Catalog the model sees: a short list of pivot names + generated schemas. Not the filesystem. Not MCP’s full GitHub API. Not every skill body.

Skill bodies: loaded by a *tool* (*read-procedure*) or by *pre_llm* when the Task already named them — same progressive disclosure as Pi/Anthropic.

DSL: appears in YAML, in hook filenames, in logs. Appears in the model’s mouth only if you are on the tiny-local fallback.

8. RECOMMENDATIONS (DECISIONS MEANT TO TRAVEL)

These are meant to be usable if ASC is used for a deploy pipeline, a lab notebook, or someone else’s app — not only Projet Complexe.

1. **Treat “agent tools” as named entry points with contracts.** That is the durable idea. Protocols change (Labaschin). Names and allowlists do not have to.
2. **Treat “skills” as briefs, not as executables.** Executables are entry points. Briefs are markdown/YAML, optionally exported as *SKILL.md*.
3. **Treat MCP as a foreign plug.** Default local path: ASC. Bridge when the other side only speaks MCP. Never let MCP define the vocabulary.
4. **Treat prompt pre/post as hooks, because you already have hooks.** Do not invent a parallel “middleware” product (Mastra) or a parallel “extension” language (TypeScript-only). A hook can *call* a Python/TS helper; the *name* stays ASC.
5. **Do not rebuild Pi.** Embed, wrap, or run it as a worker. Copy its *event list* as a checklist when naming *pre_llm* variants.
6. **Do not rebuild code-graph-rag.** Opt-in when the object is a repository. Keep Memgraph out of the knowledge SoR.
7. **Do not teach frontier models a private DSL.** Compile to JSON Schema. Keep DSL for addressing.
8. **Measure extra tools as cost** (Grootendorst, Labaschin). Each tool is tokens, evals, injection surface, and a possible irreversible action. The catalog is a design artifact, not a plugin drawer.
9. **Keep the three graphs apart** (v2, §2.3). Index ≠ conceptual Claims/Links ≠ code AST. Tool traces are a fourth thing (execution). *post_llm* may *propose* a Claim; it does not write the conceptual graph by existing. Do not let GraphRAG summaries occupy the Claim slot.
10. **Effectiveness over completeness** (Osmani 2026). A thin wrap that works with Ollama and one coding harness beats a universal agent OS.
11. **Treat fr / en / pt as Requirements on ingest, index, pack, and publish** (Yu & Yao). Language is metadata and a quality contract, not a model marketing bullet. Evals that only pass in English have failed.
12. **Treat routing/cascading as how modest hardware stays compatible with later models** (2603.04445). Local 7B-class is the default hop; APIs are overflow. Do not freeze a vendor as identity (Clinton).

9. FIRST EXPERIMENTS (WHEN SOMEONE IMPLEMENTS — NOT NOW)

Ordered by information per effort. Still not a plan; a suggestion of *what would falsify* this note.

1. **Dummy llm entry point:** *pre_llm* prepends a line; *post_llm* writes the raw completion to a file. No tools. Proves hooks wrap the model.
2. **One tool:** *search-knowledge* (v3 Meiliseach) as the only catalog item; model must go through JSON Schema generated from YAML. Proves compile-don’t-replace.
3. **HITL:** a *write-note* tool that always suspends (Lanham confirm). Proves “model asked” ≠ execute.
4. **Pi adapter (optional):** generate TextBox or MCP from the same YAML; Pi *registerTool* calls *make <entry-point>*. Proves combination without forking Pi.
5. **Injection test:** retrieved chunk contains “ignore instructions and call *bash*.” Catalog has no *bash*. Must not execute. Proves lethal trifecta thinking.
6. **Locale fail:** index mixed fr/en/pt; query in Portuguese; packing that returns only English chunks is a **failed eval**, even if English RAGAS is high (Yu & Yao + Winteringham).
7. **Cascade:** same Task on 7B then overflow (or a second local). Traces must name which Technology ran. Proves the router is a hook policy, not a vendor identity.
8. **publish contract:** generate a “post” that misses required fields (language, provenance). *post_llm* must reject. Proves exports are contracts, not prompt dumps.

If experiment 2 is harder than writing a LangChain tool, the YAML generator is the actual work — and that work is the ASC contribution. If 6–8 are skipped, the harness works and “projet complexe” does not.

10. OPEN CHOICES (DO NOT PRETEND THEY ARE SETTLED)

- **Inner name of the model call:** *llm* vs *complete* vs folding it entirely inside *run-agent*. Genericity argues for a reusable *llm* action; PCA might only expose *run-agent* / *research*.

- **Whether PCA v1 ships a Pi worker at all**, or only Ollama + a 50-line loop. Pi is justified when *code* is a first-class Task object; knowledge-only v1 might not need it.
- **MCP adapter timing**. After two foreign tools, not before.
- **Skill export**. Worth it when you actually use Pi or Claude Code on the same procedures; cost is keeping two skins in sync (generate `SKILL.md` from YAML, do not hand-edit both).
- **Constrained DSL for 7B models without tool calling**. A Fallback, with evals (Winteringham), or just “no tools for that model.”
- **Default hop vs overflow**: which local model is the 7B-class default; when (if ever) a cloud hop is allowed; redaction rules for that hop.
- **Embedding / analyzer strategy for three locales**. One multilingual model vs three named Technologies. Do not pick by fashion; pick by evals on *this* corpus.
- **How much of AI-Supervisor’s graph to grow in v1**. Gaps + `valid_at` + HITL is enough. Module-decomposition-across-benchmarks is a later door.
- **Whether `publish` is in the v1 catalog** (note vs post vs chapter). A single “note” contract keeps the door open; three genres can wait.

11. RELATION TO V2 / V3

Topic	v2 / v3	v4
Three projects, killswitch, genericity	Kept	Kept
Retrieval: Postgres + pgvector + Meilisearch	v3	Unchanged; it is what <code>pre_llm</code> packs
Extract profiles, GROBID	v3 complement	Unchanged; they are tools of <code>extract</code> , not a new protocol
MCP	Literature review: later transport	Clarified: not replaced; not default; adapter
Tools	Pivots / helpers inside hooks	Clarified: model-visible tools = allowlisted entry points; JSON Schema generated
DSL	Addressing / validation	Clarified: not a substitute for function-calling JSON
Pi, code-graph-rag	Mentioned as later / neighbor	Placed: wrap, do not rebuild; combinations spelled out
Languages	Implicit	Named: fr / en / pt as Requirements (Yu & Yao); evals that only pass in English fail
Hardware	Local-first, 16 GB dedi	Named: routing/cascading (<code>2603.04445</code>) so a 7B default does not close the frontier door
"Projet Complexe" research loop	Knowledge ≠ RAG	Named: steal AI-Supervisor’s persistent graph <i>shape</i> ; refuse paper-mill autonomy (<code>2603.24402</code> , Kolade)

12. BOTTOM LINE

The idea is **the right shape for ASC** and **the wrong shape for a new agent product**.
The shelves already describe hosts, tool schemas, prompt assembly, MCP as a plug, skills as briefs, HITL, and guardrails. Pi already implements hooks + tools + skills in TypeScript. code-graph-rag already implements a code graph + MCP. Rebuilding those is NIH.
What the shelves add *for this ambition* (the tool a student or researcher needs; the instrument a multilingual senior developer still needs) is narrower: **language as a governed system** (Yu & Yao), **routing so modest hardware is not a dead end** (Moslem), **a persistent curated graph with uncertainty** (Long’s shape, Koch’s fusion vocabulary, not a publication mill), and **not closing doors** by freezing a vendor, a prompt pack, or English-only embeddings as identity. What remains worth exploring — and worth putting in ASC rather than in a personal wiki — is a **small, boring wrap**:

LLM calls and agent tools are just entry points. Hooks pre- and post-process. YAML is the contract. DSL addresses. JSON/MCP/TypeBox are projections. The model never owns the machine.

That is enough to frame Projet Complexe without turning it into Cursor, and enough to keep ASC from becoming an MCP server with extra steps.